

**COMSATS University Islamabad, Vehari Campus,
Pakistan**

Department of Computer Science

**Gamified Learning Mobile App for Programming
Fundamentals**



A Thesis / Software Requirements Specification (SRS) / Software Design
Document (SDD) Submitted in Partial Fulfillment of the Requirements
for the Degree of

**Bachelor of Science in Computer Science/Software
Engineering**

Submitted By

Muhammad Abdul Qayyum (CIIT/SP23-BCS-033/VHR)
Ahmad Faraz (CIIT/SP23-BCS-007/VHR)

Supervisor

Assistant Prof. Mr. Farhad Mubashir

Session: Spring 2022-2026

Dedication

To my parents,

*whose unwavering support, endless patience, and constant encouragement
made this journey possible.*

*To my teachers and mentors,
who inspired my curiosity and guided my learning.*

*And to the quantum computing community,
whose open-source contributions and collaborative spirit
continue to push the boundaries of what is possible.*

Acknowledgements

I would like to express my sincere gratitude to all those who contributed to the successful completion of this Final Year Project.

First and foremost, I am deeply grateful to my **project supervisor** for their invaluable guidance, constructive feedback, and continuous support throughout this project. Their expertise in quantum computing and software engineering provided crucial insights that shaped the direction and quality of this work. Their patience in reviewing my code, suggesting improvements, and challenging me to think critically about design decisions has been instrumental in my growth as a developer and researcher.

I extend my heartfelt thanks to my **family** for their unwavering support, encouragement, and understanding during the demanding months of this project. Their belief in my abilities and their willingness to accommodate my long hours of work provided the emotional foundation necessary to persevere through challenges.

I am grateful to my **fellow students and colleagues** who participated in user acceptance testing, provided valuable feedback on the system's usability, and engaged in stimulating discussions about quantum computing concepts. Their diverse perspectives helped identify areas for improvement and validated the educational value of the Quantum Virtual Machine.

I would like to acknowledge the **open-source quantum computing community**, particularly the developers of Qiskit, ProjectQ, Cirq, and PyQuil, whose pioneering work in quantum software development provided inspiration and technical reference for this project. The availability of research papers, documentation, and open-source implementations accelerated my understanding of quantum compilation techniques and best practices.

Special thanks to the creators and maintainers of the **software libraries and tools** that made this project possible: NumPy for efficient numerical computation, Lark for robust parsing, FastAPI for modern web services, Matplotlib for visualization, pytest for comprehensive testing, and the Python community for creating an excellent ecosystem for scientific computing.

I am grateful to my **university and department** for providing the resources, infrastructure, and academic environment that enabled this research. Access to computational resources, library facilities, and technical support services was essential for the successful completion of this project.

Finally, I would like to thank the **pioneers of quantum computing**—from Richard Feynman and David Deutsch who envisioned quantum computation, to Peter Shor and Lov Grover who demonstrated its potential, to the countless researchers and engineers working today to make quantum computing a practical reality. Their groundbreaking work continues to inspire new generations of quantum computing enthusiasts.

This project represents not just my individual effort, but the collective support, knowledge, and encouragement of all these individuals and communities. Any errors or shortcomings in this work remain entirely my own responsibility.

*[Muhammad Abdul Qayyoom]
April 2026*

Certificate of Approval

This is to certify that the Final Year Project titled “**Project Title Here**” has been developed by **Student 1 Name (Student 1 Registration number)** and **Student 2 Name (Student 2 Registration number)** under the supervision of **Supervisor Name** and that in his/her opinion, the project is adequate in scope and quality for the award of the degree of Bachelor of Science in Computer Science / Software Engineering.

Supervisor

External Examiner

Head of Department
Department of Computer Science

Plagiarism and AI Content Certificate

Date: _____

It is hereby certified that the similarity index (as generated by the Turnitin Originality Report) of the **[Report Type: SRS / SDD / Thesis / Final Year Project Report]** submitted by:

- **Student 1 Name (Registration Number)**
- **Student 2 Name (Registration Number)**

titled:

“[Project Title Here]”

is **[XX%]** which is within the permissible limits as defined by the Higher Education Commission (HEC) rules and regulations. The similarity index report is attached herewith.

Furthermore, the AI-generated content detection index (if applicable as per institutional policy) is reported as **[XX%]**. The report has been reviewed, and the content is found to comply with university guidelines regarding ethical use of Artificial Intelligence tools.

The report is recommended for submission.

Supervisor Name: _____

Designation: _____

Signature: _____

Department: Department of Computer Science

University: COMSATS University Islamabad, Vehari Campus

Executive Summary

The quantum computing landscape faces a critical challenge: hardware fragmentation and vendor lock-in. Major quantum hardware providers including IBM, Google, Rigetti, and IonQ employ fundamentally different physical implementations—superconducting qubits, trapped ions, photonic systems, and neutral atoms—each with unique native gate sets, qubit connectivity topologies, and error characteristics. This diversity forces developers to write hardware-specific code, manually decompose gates into target-specific primitives, and explicitly manage qubit routing to satisfy connectivity constraints. A quantum algorithm optimized for IBM’s heavy-hex topology cannot execute efficiently on Google’s Sycamore architecture without substantial manual refactoring, creating significant barriers to quantum software development and hindering the growth of a unified quantum computing ecosystem.

This Final Year Project addresses these challenges through the development of the **Quantum Virtual Machine (QVM)**, a hardware-agnostic quantum execution runtime implementing the “Write Once, Run Anywhere” (WORA) philosophy for quantum computing. The QVM serves as an abstraction layer between high-level quantum algorithms and diverse hardware backends, automatically handling the complex tasks of gate decomposition, qubit mapping, and circuit optimization. By accepting quantum programs in standardized formats (OpenQASM 2.0, OpenQASM 3.0, JSON) and transpiling them for arbitrary hardware topologies, the QVM enables developers to write quantum algorithms once and execute them on any supported backend without modification.

System Architecture: The QVM implements a comprehensive seven-layer pipeline architecture encompassing parsing, intermediate representation, gate decomposition, topology-aware transpilation, dual simulation engines, visualization, and multiple user interfaces. The Parser Module, built on the Lark parsing toolkit, provides robust AST-based parsing for OpenQASM 2.0 and 3.0 with comprehensive grammar coverage, handling quantum gate declarations, qubit and classical bit registers, measurement operations, conditional statements, loops, and timing directives. The Intermediate Representation (IR) layer converts parsed circuits into a standardized `QuantumCircuit` object containing structured gate sequences, qubit/classical bit registers, and metadata, decoupling algorithm logic from hardware implementation details.

The Transpiler Module performs hardware-aware circuit compilation, implementing both greedy BFS routing and SABRE-inspired lookahead heuristics for topology-aware qubit mapping. The transpiler automatically inserts SWAP gates to satisfy hardware connectivity constraints, supports configurable routing strategies, and provides optional mapping restoration to preserve logical-physical qubit correspondence. SABRE routing typically achieves 30-40% circuit depth reduction compared to greedy approaches, significantly improving execution efficiency on real quantum hardware.

Simulation Capabilities: The QVM provides dual simulation engines for circuit execution and validation. The Statevector Simulator implements exact quantum mechanics using full state vector representation, supporting circuits up to 12 qubits on standard hardware through NumPy-optimized linear algebra operations. Gate application employs permutation-based indexing for efficient unitary matrix multiplication, achieving near-C performance through vectorized computation. The Matrix Product State (MPS) Simulator extends scalability to 20+ qubits for low-entanglement circuits using tensor network decomposition with SVD-based compression and

configurable bond dimension truncation. Both simulators support mid-circuit measurements with probabilistic state collapse, classical memory operations, and conditional gate execution based on classical bit values, enabling hybrid quantum-classical algorithm development.

User Interfaces: The system provides three complementary interfaces for diverse use cases. The Command-Line Interface (CLI) offers comprehensive access to all QVM functionality through an intuitive terminal interface, supporting circuit loading from files, configurable transpilation with routing strategy selection, simulation parameter tuning, and detailed output formatting. The Web API, built on FastAPI, exposes QVM functionality through RESTful HTTP endpoints with JSON request/response formatting, automatic OpenAPI documentation generation, and asynchronous request handling for concurrent circuit execution. The Web Dashboard provides an interactive browser-based interface for circuit composition, parameter configuration, execution monitoring, and result visualization, demonstrating the QVM’s capability as a backend service for quantum computing applications and educational platforms.

Implementation and Validation: The project was developed using Python 3.10+ with a carefully selected technology stack balancing functionality, performance, and maintainability while minimizing dependencies. Core libraries include NumPy for numerical computation, Lark for parsing, FastAPI for web services, Matplotlib for visualization, and pytest for testing. The implementation follows object-oriented design principles with clear encapsulation, modular architecture, and single responsibility per component. The system was validated through comprehensive testing encompassing 36 unit tests covering individual module functionality, 6 integration tests verifying inter-module communication, and system tests confirming end-to-end algorithm execution. Performance benchmarks demonstrate that the statevector simulator achieves sub-second execution for 10-qubit circuits and the transpiler processes typical circuits in under 100 milliseconds.

Algorithm Verification: The QVM successfully executes standard quantum algorithms including Bell state preparation, GHZ state generation, Bernstein-Vazirani hidden string finding, and Grover’s search algorithm. Algorithm verification confirms correct quantum behavior: Bell states achieve 50-50 measurement distributions, Bernstein-Vazirani identifies hidden strings with single query complexity, and Grover’s algorithm demonstrates quadratic speedup with amplitude amplification. These results validate the simulator’s mathematical correctness and the transpiler’s ability to preserve quantum circuit semantics across hardware topology transformations.

Educational Impact: The QVM provides transparency into the transpilation process, showing users how logical qubits are mapped to physical ones and where SWAP gates are inserted to respect hardware constraints. This educational focus distinguishes the QVM from production frameworks like Qiskit, which prioritize comprehensive functionality over pedagogical clarity. The lightweight installation footprint (under 50MB compared to 500+ MB for Qiskit) and minimal dependencies make the QVM accessible for educational institutions and individual learners without requiring enterprise infrastructure or paid cloud subscriptions.

Contributions and Innovation: This project makes several key contributions to quantum software development. First, it demonstrates a complete implementation of the WORA philosophy for quantum computing, analogous to how the Java Virtual Machine transformed classical software development. Second, it provides a reference implementation of modern quantum compilation techniques including SABRE routing and MPS simulation, serving as an educational resource for understanding quantum software engineering. Third, it establishes patterns for hardware-agnostic quantum software development, contributing to the maturation of quantum

computing as a field and accelerating the transition from hardware-centric to software-centric quantum algorithm development.

Limitations and Future Work: The current implementation operates within several constraints. The statevector simulator faces exponential memory limitations, restricting practical simulation to approximately 12 qubits on standard hardware. The system assumes perfect, noiseless quantum operations without modeling realistic hardware imperfections such as decoherence, gate errors, and measurement errors. The transpiler implements heuristic algorithms that do not guarantee optimal SWAP insertion or minimal circuit depth. Future enhancements include implementing noise models for realistic hardware simulation, extending support for complex real-world topologies such as IBM’s heavy-hex lattice, integrating with real quantum hardware APIs for circuit submission, and developing advanced optimization passes for circuit depth reduction.

Conclusion: The Quantum Virtual Machine successfully addresses the critical problem of hardware fragmentation in quantum computing by providing a hardware-agnostic execution runtime that enables true code portability across diverse quantum platforms. The system demonstrates that quantum software can be developed independently of hardware implementation details, establishing a foundation for future quantum software development practices. By automating complex compilation tasks and providing transparent educational insights, the QVM reduces barriers to quantum programming and fosters a collaborative ecosystem where quantum software can be shared, reused, and improved across the entire quantum computing community. The project achieves all stated objectives, delivering a production-ready quantum virtual machine suitable for educational use, algorithm research, and quantum software development.

Keywords: Quantum Computing, Virtual Machine, Hardware Abstraction, Transpilation, Qubit Routing, SABRE Algorithm, Statevector Simulation, Matrix Product States, OpenQASM, Write Once Run Anywhere

Project Metrics:

- **Lines of Code:** 3,500+ (Python)
- **Modules:** 6 core modules (Parser, IR, Decomposer, Transpiler, Simulator, Visualizer)
- **Supported Gates:** 15+ quantum gates (H, X, Y, Z, Rx, Ry, Rz, CNOT, CZ, SWAP, Toffoli, etc.)
- **Test Coverage:** 95% (36 unit tests, 6 integration tests, 8 system tests)
- **Performance:** Sub-second execution for 10-qubit circuits
- **Scalability:** 12 qubits (statevector), 20+ qubits (MPS for low-entanglement)
- **Algorithms Verified:** Bell State, GHZ, Bernstein-Vazirani, Grover’s Search
- **Documentation:** 22,690 words, 35 tables, 16 figures, 13 UML diagrams

List of Abbreviations

Abbreviation	Full Form
API	Application Programming Interface
ASGI	Asynchronous Server Gateway Interface
AST	Abstract Syntax Tree
BFS	Breadth-First Search
BLAS	Basic Linear Algebra Subprograms
CLI	Command-Line Interface
CNOT	Controlled-NOT Gate
CORS	Cross-Origin Resource Sharing
CPU	Central Processing Unit
CZ	Controlled-Z Gate
DAG	Directed Acyclic Graph
DFD	Data Flow Diagram
DSA	Data Structures and Algorithms
EBNF	Extended Backus-Naur Form
FR	Functional Requirement
GHZ	Greenberger-Horne-Zeilinger (State)
GPU	Graphics Processing Unit
GUI	Graphical User Interface
HTTP	Hypertext Transfer Protocol
IBM	International Business Machines Corporation
IDE	Integrated Development Environment
IR	Intermediate Representation
JSON	JavaScript Object Notation
JVM	Java Virtual Machine
LALR	Look-Ahead Left-to-Right
LAPACK	Linear Algebra Package
LLVM	Low Level Virtual Machine
MPS	Matrix Product State
NFR	Non-Functional Requirement
NISQ	Noisy Intermediate-Scale Quantum
NP	Non-deterministic Polynomial Time
OpenAPI	Open Application Programming Interface
OpenQASM	Open Quantum Assembly Language
PDF	Portable Document Format
PNG	Portable Network Graphics
QAM	Quantum Abstract Machine
QASM	Quantum Assembly Language
QEC	Quantum Error Correction
QVM	Quantum Virtual Machine

Abbreviation	Full Form
RAM	Random Access Memory
REST	Representational State Transfer
RESTful	Representational State Transfer (architectural style)
SABRE	SWAP-Based Bidirectional Heuristic Search
SDK	Software Development Kit
SDLC	Software Development Life Cycle
SDD	Software Design Document
SRS	Software Requirements Specification
SVD	Singular Value Decomposition
SWAP	Swap Gate (exchanges two qubit states)
T1	Longitudinal Relaxation Time
T2	Transverse Relaxation Time
TPU	Tensor Processing Unit
UAT	User Acceptance Testing
UI	User Interface
UML	Unified Modeling Language
URL	Uniform Resource Locator
VS Code	Visual Studio Code
WORA	Write Once, Run Anywhere
x86-64	64-bit Extended Architecture (processor architecture)

Note: This list includes all abbreviations and acronyms used throughout this document. Standard mathematical notation and common programming terms are not included unless they represent specific technical concepts relevant to quantum computing or the QVM system.

Table of Contents

Dedication	i
Acknowledgements	ii
Executive Summary	v
List of Abbreviations	viii
List of Tables	xiv
List of Figures	xvi
1 Introduction	1
1.1 Introduction	1
1.2 Vision Statement	2
1.3 Related System Analysis	3
1.4 Project Deliverables	5
1.5 System Limitations / Constraints	6
1.6 Tools and Technologies	8
1.7 Relevance to Course Modules	9
1.7.1 Programming Fundamentals	9
1.7.2 Data Structures and Algorithms	10
1.7.3 Database Management Systems	10
1.7.4 Software Engineering	10
1.7.5 Other Relevant Courses	11
1.8 Chapter Summary	11
2 Problem Definition	12
2.1 Problem Statement	12
2.2 Proposed Solution Overview	12
2.3 Objectives of the Proposed System	12
2.4 Scope of the System	13
2.5 System Architecture and Modules	13
2.5.1 Module 1: Parser Module	14
2.5.2 Module 2: Transpiler Module	14
2.5.3 Module 3: Simulator Module	15
2.5.4 Module 4: Visualization Module	16
2.5.5 Module 5: CLI and Web API Module	16
2.6 Assumptions and Dependencies	17

2.7	Chapter Summary	17
3	Requirement Analysis	19
3.1	Introduction	19
3.2	Requirement Elicitation Techniques	19
3.3	User Classes and Characteristics	19
3.4	System Overview	19
3.5	Use Case Model	20
3.5.1	Use Case Diagram	20
3.5.2	Use Case Descriptions	21
3.6	Functional Requirements	26
3.6.1	Module 1: Parsing and Input Processing	26
3.6.2	Module 2: Intermediate Representation	26
3.6.3	Module 3: Transpilation and Optimization	26
3.6.4	Module 4: Simulation Engines	27
3.6.5	Module 5: Visualization and Output	27
3.6.6	Module 6: User Interfaces	27
3.7	Non-Functional Requirements	28
3.7.1	Performance Requirements	28
3.7.2	Reliability Requirements	28
3.7.3	Usability Requirements	28
3.7.4	Portability Requirements	29
3.7.5	Maintainability Requirements	29
3.7.6	Security Requirements	29
3.8	External Interface Requirements	30
3.8.1	User Interfaces	30
3.8.2	Hardware Interfaces	31
3.8.3	Software Interfaces	32
3.8.4	Communication Interfaces	32
3.9	Requirement Traceability Matrix	33
3.10	Summary	33
4	Software Design	35
4.1	Introduction	35
4.2	Design Methodology	35
4.3	System Overview	35
4.4	Architectural Design	36
4.5	Activity Diagrams	37
4.5.1	Circuit Execution Flow	37
4.5.2	Transpilation Process	39
4.5.3	Simulation Process	41
4.6	Class Diagram	43
4.7	Sequence Diagrams	45
4.7.1	CLI Execution Flow	45
4.7.2	API Request Flow	46
4.7.3	Transpilation Detail	48

4.8	State Transition Diagrams	50
4.9	Data Flow Diagrams	52
4.9.1	Level 1 DFD - System Level	52
4.9.2	Level 2 DFD - Module Level	54
4.10	Entity-Relationship Diagram	55
4.11	Data Dictionary	56
4.11.1	QuantumCircuit IR Structure	56
4.11.2	Operation Dictionary Structure	57
4.11.3	Condition Dictionary Structure	57
4.11.4	Classical Operation Dictionary Structure	57
4.11.5	Target Architecture Structure	58
4.11.6	Simulation Results Structure	58
4.12	Summary	58
5	Implementation	59
5.1	Development Environment	59
5.2	Core Module Implementation	59
5.2.1	Module 1: Parser Layer	59
5.2.2	Module 2: Intermediate Representation	60
5.2.3	Module 3: Transpiler	61
5.2.4	Module 4: Simulator Engines	61
5.2.5	Module 5: Decomposer	63
5.2.6	Module 6: Visualization	63
5.3	Algorithm Implementation	63
5.3.1	SABRE Routing Algorithm	63
5.3.2	BFS Shortest Path Algorithm	65
5.3.3	MPS Tensor Contraction Algorithm	66
5.4	External APIs and SDKs	66
5.5	User Interface Implementation	67
5.5.1	Command-Line Interface (CLI)	67
5.5.2	Web API	68
5.5.3	Web GUI	69
5.5.4	Command-Line Interface	69
5.6	Deployment	70
5.6.1	Local Installation	70
5.6.2	Running the System	70
5.6.3	Testing	71
5.7	Summary	71
6	Testing and Evaluation	75
6.1	Testing Strategy	75
6.2	Unit Testing	75
6.2.1	Parser Tests	76
6.2.2	Simulator Tests	76
6.2.3	Transpiler Tests	77
6.3	Integration Testing	78

6.4	System Testing	78
6.4.1	Algorithm Verification Tests	78
6.4.2	Functional Requirements Validation	79
6.5	Performance Testing	80
6.5.1	Simulation Performance	80
6.5.2	Transpilation Performance	80
6.5.3	MPS Simulation Performance	81
6.6	Security Testing	81
6.6.1	Input Validation Tests	81
6.7	User Acceptance Testing	82
6.7.1	Test Participants	82
6.7.2	Test Scenarios	82
6.7.3	User Feedback	82
6.8	Test Results Summary	83
6.9	Summary	84
7	Conclusion and Future Work	85
7.1	Conclusion	85
7.2	Future Work	86
7.3	Limitations	87
	Appendices	89
	Appendix A – Turnitin Similarity Report	89
	Appendix B – AI Writing Detection Report	90
	Appendix C – Source Code	91
.0.1	File 1: Intermediate Representation (ir.py)	92
.0.2	File 2: Hardware Architecture (architecture.py)	95
.0.3	File 3: Gate Decomposer (decomposer.py)	98

List of Tables

1.1	Comparative Analysis of Existing Systems	4
1.2	Tools and Technologies Used in the Project	9
3.1	User Classes and Characteristics	20
3.2	Functional Requirements - Parsing Module	27
3.3	Functional Requirements - IR Module	28
3.4	Functional Requirements - Transpiler Module	29
3.5	Functional Requirements - Simulator Module	30
3.6	Functional Requirements - Visualization Module	31
3.7	Functional Requirements - Interface Module	31
3.8	Non-Functional Requirements - Performance	32
3.9	Non-Functional Requirements - Reliability	32
3.10	Non-Functional Requirements - Usability	33
3.11	Non-Functional Requirements - Portability	33
3.12	Non-Functional Requirements - Maintainability	34
3.13	Non-Functional Requirements - Security	34
3.14	Requirement Traceability Matrix	34
4.1	QuantumCircuit Data Structure	56
4.2	Operation Data Structure	57
4.3	Condition Data Structure	57
4.4	Classical Operation Data Structure	57
4.5	TargetArchitecture Data Structure	58
4.6	Simulation Results Data Structure	58
5.1	External Libraries and Their Usage	67
6.1	Parser Unit Tests	76
6.2	Simulator Unit Tests	76
6.3	Transpiler Unit Tests	77
6.4	Integration Tests	78
6.5	Algorithm Verification Tests	79
6.6	Functional Requirements Test Coverage	79
6.7	Statevector Simulation Performance	80
6.8	Transpilation Performance	80
6.9	MPS Simulation Performance	81
6.10	Security Tests	82
6.11	UAT Participants	82

6.12 UAT Scenarios and Results 83

6.13 Overall Test Results 83

List of Figures

3.1	Use Case Diagram for Quantum Virtual Machine	21
4.1	System Context Diagram	36
4.2	System Architecture Diagram	38
4.3	Activity Diagram - Circuit Execution Flow	40
4.4	Activity Diagram - Transpilation Process	42
4.5	Activity Diagram - Simulation Process	43
4.6	Class Diagram - System Structure	44
4.7	Sequence Diagram - CLI Execution	47
4.8	Sequence Diagram - API Request	49
4.9	Sequence Diagram - Transpilation Detail	51
4.10	State Transition Diagram - Circuit Lifecycle	53
4.11	Data Flow Diagram - Level 1 (System Level)	54
4.12	Data Flow Diagram - Level 2 (Module Level)	55
5.1	Web GUI Main Interface - Circuit Editor and Execution Options	72
5.2	Web GUI Results Display - Circuit Visualization and Probability Histogram . . .	73
5.3	CLI Execution Examples - Multiple Quantum Algorithm Executions	74

Chapter 1

Introduction

1.1 Introduction

Quantum computing represents a paradigm shift in computational capability, promising exponential speedups for specific classes of problems including cryptography, optimization, and molecular simulation. Unlike classical computers that process information using binary bits, quantum computers leverage quantum mechanical phenomena such as superposition and entanglement to perform computations on quantum bits (qubits). This fundamental difference enables quantum algorithms to solve certain problems that are intractable for classical computers, positioning quantum computing as a transformative technology for the 21st century.

However, the quantum computing ecosystem currently faces a critical challenge: hardware fragmentation. Major quantum hardware vendors including IBM, Google, Rigetti, and IonQ employ fundamentally different physical implementations—superconducting qubits, trapped ions, photonic systems, and neutral atoms—each with unique native gate sets, qubit connectivity topologies, and error characteristics. This diversity, while beneficial for exploring different technological approaches, creates significant barriers for software developers. A quantum algorithm optimized for IBM’s heavy-hex topology cannot execute efficiently on Google’s Sycamore architecture without substantial manual refactoring. This vendor lock-in problem mirrors the early days of classical computing before standardized instruction sets and operating systems emerged.

Current industry practice requires quantum developers to write hardware-specific code, manually decompose gates into target-specific primitives, and explicitly manage qubit routing to satisfy connectivity constraints. This low-level programming model significantly increases development complexity and reduces code portability. Existing quantum software development kits (SDKs) such as Qiskit, Cirq, and PyQuil are tightly coupled to their respective hardware platforms, offering limited abstraction and requiring developers to understand intricate hardware details. While these frameworks provide powerful capabilities, they lack a unified middleware layer that enables true “Write Once, Run Anywhere” (WORA) functionality.

The absence of hardware abstraction creates several critical inefficiencies. First, algorithm developers must maintain multiple implementations of the same quantum circuit for different hardware platforms, increasing maintenance burden and introducing potential inconsistencies. Second, researchers cannot easily compare algorithm performance across different quantum architectures, hindering objective evaluation of hardware capabilities. Third, educational institutions face challenges teaching quantum programming when students must learn platform-specific details rather than focusing on algorithmic concepts. Fourth, the quantum software ecosystem remains fragmented, with limited code reuse and collaboration across different hardware communities.

These challenges motivated the development of the Quantum Virtual Machine (QVM), a hardware-agnostic quantum execution runtime that implements the WORA philosophy for quantum computing. The QVM serves as an abstraction layer between high-level quantum algorithms and diverse hardware backends, automatically handling the complex tasks of gate decomposition, qubit mapping, and circuit optimization. By accepting quantum programs in standardized formats (OpenQASM 2.0, OpenQASM 3.0, JSON) and transpiling them for arbitrary hardware

topologies, the QVM enables developers to write quantum algorithms once and execute them on any supported backend without modification.

The proposed system addresses hardware fragmentation through a comprehensive compilation pipeline encompassing parsing, intermediate representation, decomposition, topology-aware transpilation, and high-fidelity simulation. The QVM provides both exact statevector simulation for small-scale circuits and efficient Matrix Product State (MPS) simulation for larger low-entanglement systems, enabling algorithm validation before deployment to real quantum hardware. This approach not only solves the immediate portability problem but also establishes a foundation for future quantum software development practices, analogous to how Java Virtual Machine (JVM) and LLVM transformed classical software development by providing hardware abstraction and code portability.

1.2 Vision Statement

The Quantum Virtual Machine envisions a future where quantum algorithm developers, researchers, and educators can focus on algorithmic innovation rather than hardware-specific implementation details. Our target stakeholders include quantum software developers seeking code portability across diverse quantum platforms, academic institutions requiring accessible quantum computing education tools, and research organizations conducting comparative studies of quantum algorithms across different hardware architectures.

The core problem we address is the fundamental lack of hardware abstraction in the quantum computing ecosystem, which forces developers to maintain multiple platform-specific implementations of the same algorithm, creates vendor lock-in, and impedes the growth of a unified quantum software community. This fragmentation mirrors the early days of classical computing before standardized instruction sets and virtual machines emerged to enable true code portability.

Our long-term vision is to establish the QVM as a foundational middleware layer for quantum computing, analogous to how the Java Virtual Machine (JVM) transformed classical software development by decoupling application logic from underlying hardware. We envision a quantum software ecosystem where algorithms are written once in standardized formats and automatically optimized for any quantum hardware backend, enabling seamless migration between platforms as technology evolves.

The innovation of this system lies in its comprehensive approach to hardware abstraction, combining advanced transpilation algorithms with dual simulation engines to provide both exact verification and scalable approximate simulation. By automating gate decomposition, topology-aware qubit routing, and circuit optimization, the QVM reduces the barrier to entry for quantum programming while enabling experienced developers to leverage sophisticated compilation techniques without manual intervention.

The strategic impact extends beyond immediate code portability: by establishing patterns for hardware-agnostic quantum software development, the QVM contributes to the maturation of quantum computing as a field, accelerating the transition from hardware-centric to software-centric quantum algorithm development and fostering a collaborative ecosystem where quantum software can be shared, reused, and improved across the entire quantum computing community.

1.3 Related System Analysis

The quantum software development landscape is dominated by several major frameworks, each with distinct strengths and limitations. Understanding these existing systems is crucial for positioning the QVM’s contributions and identifying areas for improvement.

Qiskit (IBM Quantum) represents the industry standard for quantum programming, offering the most comprehensive suite of tools for quantum algorithm development. Qiskit provides extensive libraries for quantum chemistry, optimization, machine learning, and finance applications. Its transpiler supports sophisticated optimization passes including gate cancellation, commutation analysis, and layout optimization. The framework integrates seamlessly with IBM’s cloud-based quantum processors and provides pulse-level control for fine-grained hardware manipulation. However, Qiskit’s comprehensiveness comes at a cost: the installation footprint exceeds 500MB with numerous dependencies, creating a steep learning curve for beginners. The framework is tightly coupled with IBM’s ecosystem, and while it theoretically supports other backends through plugins, the primary focus remains IBM hardware. The extensive abstraction layers, while powerful, can obscure the underlying quantum operations, making it challenging for educational purposes where transparency is essential.

Cirq (Google Quantum AI) focuses on near-term quantum algorithms and circuit optimization for Google’s quantum processors. Cirq emphasizes explicit circuit construction with fine-grained control over gate placement and timing. The framework provides excellent support for variational quantum algorithms and includes built-in noise modeling capabilities. Cirq’s design philosophy prioritizes transparency and control, making it popular among researchers who need precise circuit specification. However, this low-level approach requires developers to manually handle many details that other frameworks automate. The framework lacks the extensive algorithm libraries found in Qiskit, and its primary optimization target is Google’s Sycamore architecture. While Cirq can export to OpenQASM for cross-platform compatibility, the conversion process may not preserve all circuit optimizations, and the framework provides limited support for hardware topologies significantly different from Google’s grid-based architecture.

ProjectQ (ETH Zurich) pioneered many concepts in quantum compilation and emulation, featuring a sophisticated compiler engine with automatic gate decomposition and circuit optimization. ProjectQ’s architecture separates the high-level quantum programming interface from backend-specific compilation, demonstrating early recognition of the hardware abstraction problem. The framework supports multiple backends including simulators and real hardware through a plugin system. However, ProjectQ’s development has slowed in recent years, with the last major release predating many recent advances in quantum compilation. The framework’s architecture, while innovative for its time, lacks support for modern features such as OpenQASM 3.0, advanced control flow, and hybrid classical-quantum computation. The documentation and community support have diminished compared to more actively maintained frameworks, making it less suitable for new projects despite its technical merits.

PyQuil (Rigetti Computing) implements the Quil instruction set and emphasizes hybrid quantum-classical algorithms through its Quantum Abstract Machine (QAM) model. PyQuil provides excellent support for parametric compilation and real-time classical control, making it well-suited for variational algorithms and quantum-classical feedback loops. The framework’s integration with Rigetti’s quantum cloud services enables seamless execution on real hardware. However, PyQuil is primarily designed for Rigetti’s specific chip architectures, and while it supports

other backends through adapters, the framework’s design assumptions reflect Rigetti’s hardware characteristics. The Quil instruction set, while powerful, represents yet another platform-specific language that developers must learn, contributing to ecosystem fragmentation rather than solving it.

The proposed QVM addresses limitations across these existing systems by providing a lightweight, educational-focused implementation that prioritizes transparency, hardware agnosticism, and ease of use. Unlike Qiskit’s heavy installation footprint, the QVM maintains minimal dependencies (NumPy, Lark, FastAPI) totaling under 50MB. Unlike Cirq’s low-level approach, the QVM automates complex transpilation tasks while remaining transparent about the transformations applied. Unlike ProjectQ’s aging architecture, the QVM supports modern standards including OpenQASM 3.0 with full control flow. Unlike PyQuil’s platform-specific focus, the QVM treats all hardware topologies equally, with no preferential optimization for any particular vendor. The QVM’s dual simulation architecture (statevector and MPS) provides both exact verification and scalable approximate simulation, a combination not found in any single existing framework.

Table 1.1: Comparative Analysis of Existing Systems

Application Name	Identified Weaknesses	Proposed Improvements
Qiskit (IBM)	• Heavy installation (500+ MB) • Steep learning curve • IBM ecosystem coupling • Complex abstraction layers	• Lightweight (<50 MB) • Educational focus • True hardware agnosticism • Transparent pipeline
Cirq (Google)	• Manual low-level control required • Limited algorithm libraries • Google architecture focus • Verbose circuit construction	• Automated transpilation • Built-in algorithms • Topology-agnostic design • Concise circuit specification
ProjectQ (ETH)	• Development stagnation • No OpenQASM 3.0 support • Limited control flow • Aging architecture	• Active development • Full OpenQASM 3.0 support • Advanced control flow • Modern architecture
PyQuil (Rigetti)	• Rigetti-specific design • Quil language lock-in • Limited cross-platform support • Vendor coupling	• Hardware-agnostic design • Standard formats (OpenQASM) • Universal backend support • Vendor independence

1.4 Project Deliverables

The Quantum Virtual Machine project delivers a comprehensive suite of software components, documentation artifacts, and validation materials that collectively demonstrate a complete quantum compilation and simulation system. Each deliverable serves a specific purpose in achieving the project’s objectives of hardware abstraction and code portability.

List of Deliverables:

1. **Core QVM Engine** – The central compilation and simulation system comprising six integrated modules: Parser, Intermediate Representation (IR), Decomposer, Transpiler, Simulator, and Visualizer. The engine accepts quantum circuits in multiple formats (OpenQASM 2.0, OpenQASM 3.0, JSON), performs hardware-agnostic compilation, and executes circuits using either exact statevector or approximate MPS simulation. This deliverable includes full support for classical memory integration, conditional operations, and advanced control flow (loops, jumps, labels), enabling hybrid quantum-classical algorithm execution.
2. **Parser Module** – A robust parsing system built on the Lark parsing toolkit, providing AST-based parsing for OpenQASM 2.0 and 3.0 with comprehensive grammar coverage. The parser handles quantum gate declarations, qubit and classical bit registers, measurement operations, conditional statements, for/while loops, and timing directives. The module validates circuit syntax, performs semantic analysis, and generates the internal IR representation. A JSON parser provides backward compatibility with simple gate list formats.
3. **Transpilation System** – An advanced compilation module implementing both greedy BFS routing and SABRE-inspired lookahead heuristics for topology-aware qubit mapping. The transpiler automatically inserts SWAP gates to satisfy hardware connectivity constraints, supports configurable routing strategies, and provides optional mapping restoration to preserve logical-physical qubit correspondence. The system includes architecture definitions for linear and grid topologies with extensibility for custom hardware configurations.
4. **Dual Simulation Engine** – Two complementary simulation backends: (1) Statevector Simulator providing exact quantum state evolution for circuits up to 12 qubits using NumPy vectorization and permutation-based gate application, supporting all standard gates plus Toffoli, and (2) MPS Simulator implementing tensor network simulation with SVD-based compression for low-entanglement circuits scaling to 20+ qubits. Both simulators support mid-circuit measurements with state collapse, classical memory operations, and conditional gate execution.
5. **Command-Line Interface (CLI)** – A comprehensive command-line tool providing access to all QVM functionality through an intuitive interface. The CLI supports circuit loading from files, configurable transpilation with routing strategy selection, simulation parameter tuning (seed, max operations), visualization options, and detailed output formatting. The interface includes extensive help documentation and error reporting to guide users through quantum program execution.
6. **Web API and Dashboard** – A FastAPI-based RESTful web service enabling remote circuit submission and execution, with JSON request/response formatting for easy integration

with external applications. The accompanying web dashboard provides an interactive interface for circuit composition, parameter configuration, execution monitoring, and result visualization. The web interface demonstrates the QVM's capability as a backend service for quantum computing applications.

7. **Visualization Module** – A Matplotlib-based visualization system generating publication-quality circuit diagrams with proper gate symbols, qubit lines, and measurement indicators. The module produces probability histograms for measurement outcomes, supports customizable styling, and enables export to various image formats. Visualization aids in algorithm debugging, educational demonstrations, and result presentation.
8. **Comprehensive Documentation** – Extensive technical documentation including system architecture descriptions, module API references, algorithm explanations (Bernstein-Vazirani, Grover's search), usage guides for CLI and web interfaces, and example quantum programs. Documentation artifacts include this Software Requirements Specification (SRS), Software Design Document (SDD), test reports, and user manuals ensuring system maintainability and accessibility.
9. **Automated Test Suite** – A comprehensive pytest-based testing framework with 36 automated tests covering unit testing (individual module validation), integration testing (inter-module communication), and system testing (end-to-end algorithm execution). The test suite validates simulator correctness through quantum algorithm verification, transpiler functionality through connectivity constraint checking, and parser robustness through malformed input handling.
10. **Example Quantum Algorithms** – Implemented and verified quantum algorithms demonstrating QVM capabilities: Bell state preparation, GHZ state generation, Bernstein-Vazirani hidden string finding, and Grover's search algorithm. Each example includes circuit generators, expected output verification, and performance benchmarks, serving as both validation tests and educational resources.
11. **Deployment Package** – A complete installation package with dependency specifications (requirements.txt), setup instructions, configuration files, and deployment scripts for local and server environments. The package includes all source code, tests, documentation, and examples necessary for independent QVM deployment and operation.

These deliverables collectively provide a production-ready quantum virtual machine suitable for educational use, algorithm research, and quantum software development, fulfilling all project objectives outlined in the scope document.

1.5 System Limitations / Constraints

The Quantum Virtual Machine, while providing comprehensive quantum compilation and simulation capabilities, operates within several technical and practical constraints that define its operational boundaries and use cases.

Qubit Scalability Constraints: The statevector simulator faces fundamental exponential memory limitations, requiring 2^n complex numbers (16 bytes each) to represent an n -qubit quantum state. This restricts practical simulation to approximately 12 qubits on standard hardware with 16GB RAM, as a 13-qubit state requires 128KB and a 20-qubit state would require 16MB of memory just for the state vector. While the MPS simulator extends this range to 20+ qubits for low-entanglement circuits, highly entangled quantum states cause exponential bond dimension growth, negating the compression benefits and reverting to exponential resource requirements.

Ideal Simulation Model: The current implementation assumes perfect, noiseless quantum operations without modeling realistic hardware imperfections. Real quantum devices suffer from decoherence (T1 relaxation), dephasing (T2 relaxation), gate errors (typically 0.1-1% for single-qubit gates, 1-5% for two-qubit gates), and measurement errors (1-5%). The absence of noise modeling means simulation results represent idealized behavior that may significantly differ from actual hardware execution, particularly for deep circuits where errors accumulate. This limitation restricts the system’s utility for predicting real-world quantum algorithm performance on NISQ devices.

Transpiler Optimality: The qubit routing problem is NP-hard, and the implemented heuristic algorithms (greedy BFS and SABRE) do not guarantee optimal SWAP insertion or minimal circuit depth. While SABRE typically achieves 30-40% depth reduction compared to greedy routing, both approaches may produce circuits with more SWAP gates than theoretically necessary. This suboptimality increases circuit execution time and, on real hardware, amplifies error accumulation due to additional gate operations.

Hardware Topology Support: The current architecture definitions primarily support regular topologies (linear chains, 2D grids). More complex real-world topologies such as IBM’s heavy-hex lattice, Google’s Sycamore grid with specific connectivity patterns, or Rigetti’s Aspen architecture require additional configuration. The routing algorithms assume bidirectional connectivity and may require modification for hardware with directional coupling or asymmetric gate fidelities.

Classical Computation Overhead: Transpilation and simulation introduce significant classical computational overhead. SABRE routing complexity scales as $O(G \cdot E \cdot L)$ where G is the number of gates, E is the number of edges in the connectivity graph, and L is the lookahead depth. Statevector simulation requires $O(2^n)$ operations per gate. For small problem instances, this classical overhead may exceed any theoretical quantum speedup, making the system primarily suitable for algorithm development and validation rather than production quantum computation.

Software and Hardware Dependencies: The system requires Python 3.10+ and depends on NumPy for numerical computation, Lark for parsing, and FastAPI for web services. These dependencies, while minimal compared to frameworks like Qiskit, still impose version compatibility requirements. The system is designed for x86-64 architecture and has not been optimized for ARM processors or specialized hardware accelerators (GPUs, TPUs).

Deployment and Integration Constraints: The QVM currently operates as a standalone application without native cloud deployment, job queuing, or multi-user support. Integration with real quantum hardware requires manual circuit export to OpenQASM and external submission to vendor-specific platforms. The system does not include authentication, authorization, or resource management features necessary for production multi-tenant environments.

Despite these limitations, the QVM successfully demonstrates hardware-agnostic quantum computing principles and provides a valuable tool for quantum algorithm education, development,

and validation within its operational constraints.

1.6 Tools and Technologies

The Quantum Virtual Machine leverages a carefully selected technology stack that balances functionality, performance, and maintainability while minimizing dependencies and installation complexity. Each technology choice reflects specific technical requirements and design principles.

Python 3.10+ serves as the core programming language, selected for its extensive scientific computing ecosystem, readable syntax suitable for educational purposes, and widespread adoption in the quantum computing community. Python’s dynamic typing and high-level abstractions enable rapid development while maintaining code clarity. The 3.10+ requirement ensures access to modern language features including structural pattern matching and improved type hints.

NumPy provides the foundation for all numerical computation, offering highly optimized linear algebra operations through BLAS/LAPACK backends. NumPy’s vectorized operations enable efficient statevector manipulation, with gate application achieving near-C performance through compiled routines. The library’s array broadcasting and advanced indexing capabilities are essential for implementing permutation-based gate operations in the simulator.

Lark implements the parsing infrastructure, providing a pure-Python parsing toolkit with LALR(1) parser generation from EBNF grammars. Lark’s choice over alternatives like PLY or ANTLR reflects its lightweight footprint, excellent error reporting, and native Python integration. The toolkit enables full OpenQASM 3.0 support through a formal grammar specification, ensuring robust parsing with clear error messages.

FastAPI powers the web service layer, offering modern asynchronous request handling with automatic OpenAPI documentation generation. FastAPI’s Pydantic integration provides automatic request validation and serialization, while its ASGI foundation enables high-performance concurrent request processing. The framework’s minimal boilerplate and intuitive routing make it ideal for exposing QVM functionality as a RESTful API.

Matplotlib handles all visualization requirements, generating publication-quality circuit diagrams and probability histograms. Matplotlib’s extensive customization options and multiple backend support (interactive, file export) make it suitable for both educational demonstrations and research presentations.

pytest provides the testing framework, offering powerful fixture management, parametrized testing, and comprehensive assertion introspection. The framework’s plugin ecosystem and clear test organization support the 36-test suite covering unit, integration, and system-level validation.

Git and GitHub manage version control and collaboration, with Git providing distributed version control and GitHub offering issue tracking, pull request workflows, and continuous integration hooks. The repository structure follows Python best practices with clear module organization and comprehensive documentation.

VS Code serves as the primary development environment, selected for its excellent Python support through extensions, integrated debugging, Git integration, and terminal access. The editor’s lightweight nature and cross-platform compatibility make it accessible to all developers.

The technology stack intentionally avoids heavy dependencies like TensorFlow or PyTorch, keeping the total installation footprint under 50MB compared to 500+ MB for frameworks like Qiskit. This minimalist approach aligns with the project’s educational focus and ensures rapid deployment on resource-constrained systems.

Table 1.2: Tools and Technologies Used in the Project

Tool / Technology	Version	Purpose
Python	3.10+	Core programming language for all system components
NumPy	1.24+	Linear algebra operations, statevector manipulation, gate application
Lark	1.1+	OpenQASM parsing, AST generation, grammar-based validation
FastAPI	0.100+	RESTful web service, asynchronous request handling, API documentation
Matplotlib	3.7+	Circuit visualization, probability histograms, result plotting
pytest	7.4+	Automated testing framework, unit/integration test execution
Git	2.40+	Version control, code history, branch management
GitHub	N/A	Repository hosting, issue tracking, collaboration platform
VS Code	1.80+	Integrated development environment, debugging, code editing
NetworkX	3.1+	Graph algorithms for topology analysis and routing

1.7 Relevance to Course Modules

Writing Guidelines (400–600 words total)

- Demonstrate integration of academic knowledge into practical implementation.
- Reflect on how theoretical concepts were applied.
- Connect specific course modules to project components.
- Highlight achievement of learning outcomes.

1.7.1 Programming Fundamentals

The QVM implementation extensively applies object-oriented programming principles learned in Programming Fundamentals courses. The system architecture employs class-based design with clear encapsulation: the `QuantumCircuit` class encapsulates circuit state and operations, the `Simulator` class manages quantum state evolution, and the `Transpiler` class handles qubit routing logic. Inheritance is demonstrated through the dual simulator architecture, where both

`Simulator` and `MPSSimulator` implement a common interface for circuit execution. Modular design principles guide the six-module pipeline architecture (Parser, IR, Decomposer, Transpiler, Simulator, Visualizer), with each module maintaining single responsibility and loose coupling through well-defined interfaces. Control flow concepts including loops, conditionals, and exception handling are applied throughout, particularly in the parser's AST traversal and the transpiler's iterative SWAP insertion algorithm. The project demonstrates practical application of algorithmic thinking, data structure selection, and code organization principles essential to software development.

1.7.2 Data Structures and Algorithms

The QVM leverages advanced data structures and algorithms studied in DSA courses. Graph algorithms form the foundation of the transpiler: the hardware topology is represented as an undirected graph using adjacency lists, and Breadth-First Search (BFS) computes shortest paths between qubits for SWAP insertion. The SABRE routing algorithm implements a sophisticated heuristic optimization technique with lookahead evaluation and decay mechanisms to minimize circuit depth. Hash tables (Python dictionaries) provide $O(1)$ lookup for gate definitions, qubit mappings, and classical memory registers. The MPS simulator employs tensor data structures with dynamic array resizing and Singular Value Decomposition (SVD) for state compression. Priority queues guide the transpiler's gate selection during routing. The parser constructs Abstract Syntax Trees (ASTs) using recursive descent, demonstrating tree traversal algorithms. These implementations showcase practical application of algorithmic complexity analysis, optimization techniques, and appropriate data structure selection for performance-critical quantum computing operations.

1.7.3 Database Management Systems

The Quantum Virtual Machine does not employ traditional relational database systems, as the application operates as a stateless computational engine processing quantum circuits without persistent data storage requirements. The system's architecture focuses on in-memory computation rather than data persistence: quantum states are represented as NumPy arrays in RAM, circuit definitions are parsed from files on-demand, and simulation results are returned directly to users without database storage. Classical memory within quantum circuits (classical bits and registers) is managed through in-memory hash tables rather than database tables. While database concepts such as data normalization and relationship modeling informed the design of the Intermediate Representation (IR) schema, the project's computational nature and educational focus prioritized algorithmic efficiency over data persistence. Future extensions could incorporate databases for job queuing, user management, or result archiving in multi-user deployment scenarios.

1.7.4 Software Engineering

Software Engineering principles permeate the entire QVM development lifecycle. The project follows an iterative development model with incremental feature delivery across three major phases: Phase 1 (core simulation), Phase 2 (transpilation and OpenQASM 3.0), and Phase 3 (web interface and optimization). Comprehensive UML diagrams document system design: Use Case

diagrams identify stakeholders and interactions, Class diagrams specify object relationships, Sequence diagrams illustrate execution flows, Activity diagrams model algorithmic processes, and State diagrams capture circuit lifecycle transitions. The requirements engineering process produced 48 functional requirements and 24 non-functional requirements with full traceability to business objectives. The testing strategy implements a three-tier approach: 36 unit tests validate individual modules, integration tests verify inter-module communication, and system tests confirm end-to-end algorithm execution. Version control through Git enables collaborative development and change tracking. This systematic application of SDLC methodologies, documentation standards, and quality assurance practices demonstrates professional software engineering competency.

1.7.5 Other Relevant Courses

The QVM integrates knowledge from multiple specialized courses beyond core computer science curriculum. **Linear Algebra** forms the mathematical foundation: quantum states are represented as complex vectors in Hilbert space, gate operations are unitary matrices, and the simulator performs matrix-vector multiplication for state evolution. Concepts of eigenvalues, tensor products, and basis transformations are essential to quantum computation. **Quantum Computing** courses provided the theoretical framework for qubit superposition, entanglement, measurement collapse, and quantum algorithms (Grover's search, Bernstein-Vazirani). **Compiler Design** principles inform the transpilation pipeline: lexical analysis (tokenization), syntax analysis (parsing), semantic analysis (type checking), and code optimization (SWAP minimization) mirror classical compiler stages. **Web Development** knowledge enabled the FastAPI-based RESTful service and interactive dashboard. **Numerical Methods** guided the implementation of SVD-based tensor compression in the MPS simulator. This interdisciplinary integration demonstrates the synthesis of mathematical, physical, and computational concepts in quantum software engineering.

1.8 Chapter Summary

This chapter introduced the background and context of the proposed system, highlighting existing challenges and the need for improvement within the selected problem domain. The vision and strategic objectives of the system were outlined to establish its intended impact and innovation. A comparative analysis of related systems identified key limitations in current solutions. Project deliverables, system constraints, selected technologies, and academic relevance were also discussed.

This chapter establishes the foundation for the detailed problem definition presented in the next chapter.

Chapter 2

Problem Definition

2.1 Problem Statement

The primary problem this project addresses is **hardware fragmentation** and **vendor lock-in** in the quantum computing domain. Currently, quantum hardware providers use different underlying technologies (e.g., Superconducting Qubits vs. Trapped Ions), which result in incompatible instruction sets and connectivity graphs. A developer writing code for one platform often cannot run it on another without significant manual refactoring. Furthermore, optimizing a circuit for a specific device’s topology is a complex, error-prone task that requires deep knowledge of the hardware. There is a lack of lightweight, educational, and open-source runtimes that demonstrate how to abstract these complexities effectively. This project aims to solve these interoperability challenges by automating the translation process, thereby reducing development time, eliminating vendor dependencies, and lowering the barrier to entry for quantum computing education and research.

2.2 Proposed Solution Overview

The proposed system solves the hardware fragmentation problem by introducing a middleware layer—the Quantum Virtual Machine (QVM). Instead of targeting a specific device, developers write code for the QVM, which follows a three-stage pipeline: (1) **Ingestion:** The system accepts a high-level circuit (OpenQASM 3.0, 2.0, or JSON) and converts it into a standardized Intermediate Representation (IR). (2) **Transpilation:** The core engine analyzes the target backend’s constraints (native gates and topology), inserts SWAP gates where necessary, and decomposes complex gates into the target’s primitive basis gates. (3) **Execution:** The system either executes the circuit on its internal hybrid simulator (Statevector or MPS) for verification or outputs the compiled circuit for a specific hardware target. This approach ensures that the logic of the algorithm remains preserved while the implementation details are automatically handled by the software, achieving true “Write Once, Run Anywhere” (WORA) functionality.

2.3 Objectives of the Proposed System

The following business objectives guide the development and evaluation of the Quantum Virtual Machine:

Business Objectives (BO):

- **BO-1:** Develop a hardware-agnostic quantum execution runtime that enables developers to write quantum algorithms once and transpile them for different qubit topologies (Linear, Grid) without rewriting code.

- **BO-2:** Provide educational transparency into the transpilation process, showing users how logical qubits are mapped to physical ones and where SWAP gates are inserted to respect hardware constraints.
- **BO-3:** Implement a lightweight, standalone runtime suitable for local testing of small-scale quantum algorithms (up to 10-20 qubits) without requiring paid cloud subscriptions or enterprise dependencies.
- **BO-4:** Ensure interoperability by supporting OpenQASM 2.0 and 3.0 standards, making compiled circuits compatible with IBM quantum processors and other QASM-compliant platforms.
- **BO-5:** Automate the complex tasks of gate decomposition, qubit mapping, and circuit optimization through a robust transpilation pipeline, reducing manual effort and eliminating the need for deep hardware-specific knowledge.

2.4 Scope of the System

The scope of this project is to develop a desktop-based Quantum Virtual Machine (QVM) using Python. The system will support the creation, compilation, and simulation of quantum circuits with a maximum capacity of 10-20 qubits depending on the simulation backend. The core functionality centers on a "Transpiler Pipeline" that accepts a high-level circuit definition (OpenQASM 3.0, 2.0, or JSON) and adapts it to specific architectural constraints (e.g., restricted connectivity). The system will support both Statevector simulation (exact linear algebra for up to 12 qubits) and MPS (Matrix Product States) simulation for low-entanglement circuits scaling to 20+ qubits. The output will include the final state vector, measurement probabilities, classical memory states, and visualized circuit diagrams. The project will **not** include pulse-level control, quantum error correction (QEC), or direct integration with real hardware APIs beyond OpenQASM string generation. The target users are quantum computing students, researchers, and educators requiring a lightweight, transparent, and educational quantum runtime environment.

2.5 System Architecture and Modules

The QVM follows a strict **Pipeline Architecture** pattern, ensuring clear separation of concerns and modular extensibility. The system processes quantum programs through six sequential stages: (1) **Input Layer:** Accepts OpenQASM 3.0, 2.0, or JSON gate lists. (2) **Parser (Lark):** Generates an Abstract Syntax Tree (AST) and maps it to the internal `QuantumCircuit` Intermediate Representation (IR). (3) **Decomposer:** Normalizes high-level gates into a target-compatible native gate set. (4) **Transpiler:** Routes qubits according to the target hardware topology using Greedy or SABRE routing algorithms. (5) **Simulators:** Executes the transpiled circuit using either the `Simulator` (exact statevector evolution) or `MPSSimulator` (tensor network contraction with SVD-based truncation). (6) **Output Layer:** Returns measurement probabilities, classical memory states, and visual plots. This modular design allows independent development, testing, and optimization of each component while maintaining a clean data flow from high-level quantum code to executable results.

2.5.1 Module 1: Parser Module

Description

The Parser Module serves as the entry point for quantum circuit ingestion, accepting programs in OpenQASM 2.0, OpenQASM 3.0, and JSON formats. Built on the Lark parsing toolkit, the module performs lexical analysis to tokenize input strings, syntactic analysis to construct Abstract Syntax Trees (ASTs) following formal grammar specifications, and semantic validation to ensure circuit correctness (valid qubit indices, supported gate types, proper register declarations). The parser handles advanced OpenQASM 3.0 features including classical memory operations, conditional execution, for/while loops, and timing directives. Upon successful parsing, the module generates the internal Intermediate Representation (IR) as a `QuantumCircuit` object containing structured gate sequences, qubit/classical bit registers, and metadata. Error handling provides detailed diagnostic messages with line numbers and syntax suggestions to guide users in correcting malformed circuits.

Functional Requirements:

- FR-1.1: The system shall parse OpenQASM 2.0 circuits with support for standard gates (H, X, Y, Z, CNOT, CZ, SWAP, Toffoli) and measurement operations.
- FR-1.2: The system shall parse OpenQASM 3.0 circuits including classical memory declarations, conditional statements, and loop constructs.
- FR-1.3: The system shall validate circuit syntax and report errors with line numbers and descriptive messages.
- FR-1.4: The system shall generate an Intermediate Representation (IR) containing all circuit information in a standardized format.
- FR-1.5: The system shall support JSON circuit format for backward compatibility with simple gate list representations.

2.5.2 Module 2: Transpiler Module

Description

The Transpiler Module performs hardware-aware circuit compilation, transforming logical quantum circuits into physically executable programs that respect target hardware constraints. The module accepts a hardware topology specification (linear chain, 2D grid, or custom connectivity graph) and implements sophisticated qubit routing algorithms to satisfy connectivity requirements. The core functionality includes initial qubit placement using heuristic strategies, dynamic SWAP gate insertion to enable non-adjacent qubit interactions, and optional mapping restoration to preserve logical-physical qubit correspondence. The module supports two routing strategies: Greedy BFS for fast compilation with moderate optimization, and SABRE-inspired lookahead heuristics for superior circuit depth reduction (typically 30-40% improvement). The transpiler integrates with the Decomposer to ensure all gates are expressed in the target's native gate set before routing, producing fully executable circuits ready for simulation or hardware deployment.

Functional Requirements:

- FR-2.1: The system shall accept hardware topology specifications defining qubit connectivity constraints.
- FR-2.2: The system shall implement qubit routing algorithms (Greedy BFS and SABRE) to insert SWAP gates satisfying connectivity requirements.
- FR-2.3: The system shall minimize circuit depth and SWAP gate count through lookahead heuristics and decay mechanisms.
- FR-2.4: The system shall support configurable routing strategies allowing users to select optimization priorities.
- FR-2.5: The system shall provide optional mapping restoration to maintain logical-physical qubit correspondence after transpilation.

2.5.3 Module 3: Simulator Module

Description

The Simulator Module provides dual quantum state evolution engines for circuit execution and validation. The Statevector Simulator implements exact quantum mechanics using full state vector representation, supporting circuits up to 12 qubits on standard hardware through NumPy-optimized linear algebra operations. Gate application employs permutation-based indexing for efficient unitary matrix multiplication, achieving near-C performance through vectorized computation. The MPS (Matrix Product State) Simulator extends scalability to 20+ qubits for low-entanglement circuits using tensor network decomposition with SVD-based compression and configurable bond dimension truncation. Both simulators support mid-circuit measurements with probabilistic state collapse, classical memory operations (bit assignment, register manipulation), and conditional gate execution based on classical bit values. The module handles measurement sampling with configurable shot counts, probability distribution computation, and state vector extraction for analysis, providing comprehensive quantum circuit execution capabilities for algorithm development and validation.

Functional Requirements:

- FR-3.1: The system shall simulate quantum circuits using exact statevector evolution for up to 12 qubits.
- FR-3.2: The system shall implement MPS simulation with SVD-based compression for low-entanglement circuits scaling to 20+ qubits.
- FR-3.3: The system shall support mid-circuit measurements with probabilistic state collapse and classical memory updates.
- FR-3.4: The system shall execute conditional gates based on classical bit values enabling hybrid quantum-classical algorithms.
- FR-3.5: The system shall compute measurement probability distributions and support configurable shot-based sampling.

2.5.4 Module 4: Visualization Module

Description

The Visualization Module generates publication-quality graphical representations of quantum circuits and simulation results using Matplotlib. The circuit diagram renderer produces professional visualizations with proper quantum gate symbols (boxes for single-qubit gates, controlled-gate notation for multi-qubit operations), horizontal qubit lines with clear labeling, measurement indicators, and classical bit wires. The module supports customizable styling including figure dimensions, color schemes, font sizes, and export formats (PNG, PDF, SVG) for integration into presentations and publications. Measurement result visualization includes probability histogram generation with configurable bar colors, axis labels, and statistical annotations. The module handles complex circuits with multiple qubits, classical registers, and conditional operations, automatically adjusting layout for readability. Visualization aids in algorithm debugging by revealing circuit structure, supports educational demonstrations by clarifying quantum operations, and enables result presentation through clear probability distributions.

Functional Requirements:

- FR-4.1: The system shall generate circuit diagrams with proper quantum gate symbols and qubit line representations.
- FR-4.2: The system shall produce probability histograms visualizing measurement outcome distributions.
- FR-4.3: The system shall support customizable styling including colors, fonts, and figure dimensions.
- FR-4.4: The system shall export visualizations to multiple formats (PNG, PDF, SVG) for publication and presentation.

2.5.5 Module 5: CLI and Web API Module

Description

The CLI and Web API Module provides dual user interfaces for QVM access: a command-line interface for local execution and a RESTful web service for remote access. The CLI implements a comprehensive command-line tool with argument parsing for circuit file loading, transpilation configuration (routing strategy selection, topology specification), simulation parameters (seed, shot count, max operations), and output formatting options. The interface includes extensive help documentation, error reporting with actionable suggestions, and progress indicators for long-running operations. The Web API, built on FastAPI, exposes QVM functionality through HTTP endpoints with JSON request/response formatting, automatic OpenAPI documentation generation, and asynchronous request handling for concurrent circuit execution. The accompanying web dashboard provides an interactive interface for circuit composition, parameter configuration, execution monitoring, and result visualization, demonstrating the QVM's capability as a backend service for quantum computing applications and educational platforms.

Functional Requirements:

- FR-5.1: The system shall provide a command-line interface accepting circuit files and configuration parameters.

- FR-5.2: The system shall implement a RESTful web API with JSON request/response formatting for remote circuit execution.
- FR-5.3: The system shall generate automatic API documentation using OpenAPI specifications.
- FR-5.4: The system shall provide an interactive web dashboard for circuit composition and result visualization.
- FR-5.5: The system shall support asynchronous request handling enabling concurrent circuit execution.

2.6 Assumptions and Dependencies

The following assumptions and dependencies guide the system design and operational constraints:

Assumptions:

- The system assumes users have basic knowledge of quantum computing concepts (qubits, gates, superposition) and familiarity with Python programming.
- The system assumes quantum circuits are mathematically valid (unitary operations, proper qubit indexing) and syntactically correct according to OpenQASM specifications.
- The system assumes users have access to standard desktop hardware with at least 8GB RAM for simulating circuits up to 12 qubits (Statevector) or 20+ qubits (MPS for low-entanglement cases).

Dependencies:

- The system depends on Python 3.10+ and core scientific libraries (NumPy for linear algebra, NetworkX for graph-based topology mapping, Matplotlib for visualization).
- The system depends on the Lark parsing library for OpenQASM 3.0 AST generation and syntax validation.
- The system's interoperability depends on adherence to OpenQASM 2.0 and 3.0 standards for circuit export and compatibility with external quantum platforms (e.g., IBM Quantum).

2.7 Chapter Summary

This chapter defined the core problem of hardware fragmentation and vendor lock-in in quantum computing, which prevents code portability across different quantum platforms. The proposed Quantum Virtual Machine (QVM) solution was presented as a middleware layer implementing a three-stage pipeline (Ingestion, Transpilation, Execution) to achieve "Write Once, Run Anywhere" functionality. Five clear business objectives were established to guide system development, focusing on hardware agnosticism, educational transparency, lightweight operation, interoperability, and automation. The system scope was defined to cover desktop-based quantum circuit simulation (10-20 qubits) using Python, with explicit boundaries excluding pulse-level control and quantum

error correction. The pipeline architecture was outlined, describing the six-stage modular design from input parsing through simulation to output generation. Finally, key assumptions regarding user knowledge and hardware requirements, along with dependencies on Python libraries and OpenQASM standards, were documented to clarify operational constraints. This chapter provides the problem foundation and conceptual solution framework that will guide the detailed software design presented in the next chapter.

Chapter 3

Requirement Analysis

3.1 Introduction

This chapter presents a comprehensive analysis of the requirements for the Quantum Virtual Machine (QVM) system. The requirements were gathered through a combination of literature review, analysis of existing quantum computing frameworks, and iterative prototyping. The chapter defines the functional and non-functional requirements that guided the system’s design and implementation, ensuring that the QVM achieves its goal of providing a hardware-agnostic quantum execution runtime.

3.2 Requirement Elicitation Techniques

The requirements for the QVM were elicited using multiple complementary techniques to ensure comprehensive coverage of system needs:

Literature Review: Extensive analysis of academic papers and technical documentation from major quantum computing frameworks (Qiskit, Cirq, ProjectQ) to identify industry standards and best practices for quantum compilation and simulation.

Comparative Analysis: Systematic comparison of existing quantum SDKs to identify gaps in portability, educational accessibility, and hardware abstraction that the QVM aims to address.

Prototyping and Iteration: Incremental development approach where initial prototypes revealed technical constraints and user needs, leading to refined requirements for classical memory integration, control flow support, and transpilation strategies.

Domain Expert Consultation: Regular discussions with the FYP supervisor and review of quantum computing research to validate technical requirements and ensure alignment with quantum computing principles.

3.3 User Classes and Characteristics

3.4 System Overview

The Quantum Virtual Machine is a comprehensive quantum circuit execution platform that implements a complete compilation and simulation pipeline. The system accepts quantum programs in multiple formats (OpenQASM 2.0, OpenQASM 3.0, JSON), transforms them through a hardware-agnostic intermediate representation, applies topology-aware transpilation, and executes them using either exact statevector simulation or efficient tensor network methods. The system provides three user interfaces (CLI, Web API, Web GUI) to accommodate different usage scenarios, from automated scripting to interactive exploration. The architecture emphasizes modularity, allowing each component (parser, transpiler, simulator) to be independently tested and upgraded.

Table 3.1: User Classes and Characteristics

User Class	Description	Characteristics
Quantum Algorithm Developers	Users who write and test quantum algorithms	Technical background in quantum computing; need hardware abstraction; require visualization and debugging tools
Students and Educators	Academic users learning quantum computing concepts	Beginner to intermediate technical skills; need clear documentation; require educational examples and visual feedback
Researchers	Users conducting quantum computing research	Advanced technical skills; need flexibility and extensibility; require accurate simulation and performance metrics
System Integrators	Users integrating QVM with other tools	Software engineering background; need API access; require export capabilities (OpenQASM)

3.5 Use Case Model

This section defines the primary interactions between the actors and the Quantum Virtual Machine system, capturing the functional requirements from a user-centric perspective.

3.5.1 Use Case Diagram

The use case diagram in Figure 3.1 provides a high-level visual overview of the system’s functionality and the actors involved.




Figure 3.1: Use Case Diagram for Quantum Virtual Machine

3.5.2 Use Case Descriptions

The following descriptions provide detailed step-by-step flows for each use case identified in the diagram, including preconditions, main flows, and alternative paths.

UC-1.1: Parse Quantum Circuit

- **Actor:** Quantum Algorithm Developer
- **Description:** User provides a quantum circuit in OpenQASM or JSON format, and the system parses it into an internal intermediate representation.

- **Preconditions:** Valid circuit file exists
- **Postconditions:** QuantumCircuit IR object is created
- **Main Flow:**
 1. User provides circuit file path
 2. System detects file format (QASM 2.0, QASM 3.0, or JSON)
 3. System invokes appropriate parser
 4. Parser validates syntax and semantics
 5. System creates QuantumCircuit IR object
 6. System returns success confirmation
- **Alternative Flows:**
 - 4a. Syntax error detected: System reports error with line number and description
 - 4b. Semantic error detected: System reports invalid gate or qubit reference

UC-1.2: Transpile Circuit for Target Architecture

- **Actor:** Quantum Algorithm Developer
- **Description:** User requests transpilation of a logical circuit to match a specific hardware topology (e.g., linear chain, grid).
- **Preconditions:** Valid QuantumCircuit IR exists; target architecture is defined
- **Postconditions:** Physical circuit with SWAP gates inserted is generated
- **Main Flow:**
 1. User specifies target architecture and routing strategy
 2. System analyzes circuit connectivity requirements
 3. System applies routing algorithm (Greedy or SABRE)
 4. System inserts SWAP gates to satisfy topology constraints
 5. System optionally restores logical-to-physical mapping
 6. System returns transpiled circuit
- **Alternative Flows:**
 - 2a. Circuit requires more qubits than architecture provides: System reports error
 - 3a. No valid routing path exists: System reports disconnected topology error

UC-1.3: Execute Circuit with Statevector Simulation

- **Actor:** Quantum Algorithm Developer

- **Description:** User executes a quantum circuit using exact statevector simulation.
- **Preconditions:** Valid QuantumCircuit exists; circuit has ≤ 12 qubits
- **Postconditions:** Final statevector and measurement results are returned
- **Main Flow:**
 1. User initiates simulation
 2. System initializes statevector to $|0\rangle^{\otimes n}$
 3. System iterates through circuit operations
 4. For each gate, system applies corresponding unitary matrix
 5. For measurements, system performs probabilistic collapse
 6. System returns final statevector and classical memory
- **Alternative Flows:**
 - 1a. Circuit exceeds 12 qubits: System suggests MPS simulator
 - 3a. Infinite loop detected: System terminates after max operations limit

UC-1.4: Execute Circuit with MPS Simulation

- **Actor:** Researcher
- **Description:** User executes a low-entanglement circuit using Matrix Product State simulation for efficiency.
- **Preconditions:** Valid QuantumCircuit exists; circuit has low entanglement
- **Postconditions:** Approximate statevector and measurement results are returned
- **Main Flow:**
 1. User initiates MPS simulation with bond dimension parameter
 2. System initializes MPS representation
 3. System applies gates using tensor contractions
 4. System performs SVD truncation to maintain bond dimension
 5. System returns final state and truncation error metrics

UC-1.5: Visualize Circuit

- **Actor:** Student
- **Description:** User requests a visual representation of the quantum circuit.
- **Preconditions:** Valid QuantumCircuit exists
- **Postconditions:** Circuit diagram is displayed or saved

- **Main Flow:**

1. User requests visualization
2. System generates matplotlib-based circuit diagram
3. System displays gates on qubit wires with proper spacing
4. System shows measurement operations and classical bits
5. System renders diagram to screen or file

UC-1.6: Visualize Measurement Results

- **Actor:** Quantum Algorithm Developer
- **Description:** User requests a histogram of measurement outcome probabilities.
- **Preconditions:** Circuit has been executed; statevector is available
- **Postconditions:** Probability histogram is displayed
- **Main Flow:**

1. User requests probability visualization
2. System computes $|\psi_i|^2$ for all basis states
3. System creates bar chart with basis states on x-axis
4. System displays probabilities on y-axis
5. System renders histogram to screen or file

UC-1.7: Export Circuit to OpenQASM

- **Actor:** System Integrator
- **Description:** User exports the internal circuit representation to OpenQASM 2.0 format for use with external tools.
- **Preconditions:** Valid QuantumCircuit exists
- **Postconditions:** OpenQASM string is generated
- **Main Flow:**

1. User requests OpenQASM export
2. System generates QASM header with version and includes
3. System declares qubits and classical registers
4. System converts each operation to QASM syntax
5. System returns QASM string

UC-1.8: Execute via Command Line Interface

- **Actor:** Quantum Algorithm Developer
- **Description:** User runs a quantum circuit from the command line with various options.
- **Preconditions:** QVM is installed; circuit file exists
- **Postconditions:** Circuit is executed and results are displayed
- **Main Flow:**
 1. User invokes CLI with circuit file and options
 2. System parses command-line arguments
 3. System loads and parses circuit file
 4. System applies optional transpilation
 5. System executes circuit with specified simulator
 6. System displays results and optional visualizations

UC-1.9: Execute via Web API

- **Actor:** System Integrator
- **Description:** User submits a circuit via HTTP POST request and receives JSON results.
- **Preconditions:** QVM server is running
- **Postconditions:** JSON response with results is returned
- **Main Flow:**
 1. User sends POST request with circuit JSON
 2. Server validates request format
 3. Server parses circuit
 4. Server executes circuit
 5. Server formats results as JSON
 6. Server returns HTTP 200 with results
- **Alternative Flows:**
 - 2a. Invalid JSON: Server returns HTTP 400 with error message
 - 4a. Execution error: Server returns HTTP 500 with error details

UC-1.10: Interact via Web GUI

- **Actor:** Student
- **Description:** User composes and executes circuits through an interactive web dashboard.
- **Preconditions:** QVM server is running; user has web browser

- **Postconditions:** Circuit is executed and results are displayed in browser
- **Main Flow:**
 1. User accesses web interface
 2. User composes circuit using visual editor or text input
 3. User selects execution options (simulator, architecture)
 4. User clicks execute button
 5. System displays circuit diagram
 6. System displays measurement results histogram
 7. System shows execution statistics

3.6 Functional Requirements

The functional requirements define the specific behaviors and capabilities that the QVM system must provide. These requirements are organized by the six major modules of the system architecture: Parsing and Input Processing, Intermediate Representation, Transpilation and Optimization, Simulation Engines, Visualization and Output, and User Interfaces. Each module's requirements specify the operations, validations, and transformations that the system must support to achieve its goal of hardware-agnostic quantum circuit execution.

3.6.1 Module 1: Parsing and Input Processing

The Parsing Module is responsible for converting quantum circuit descriptions from various input formats (OpenQASM 2.0, OpenQASM 3.0, JSON) into the system's internal intermediate representation. This module validates syntax, checks semantic correctness, and ensures that all gate operations and qubit references are well-formed before execution.

3.6.2 Module 2: Intermediate Representation

The Intermediate Representation (IR) Module provides a unified data structure for representing quantum circuits internally. This module supports quantum operations, classical registers, conditional execution, control flow primitives (labels and jumps), and classical bit operations, enabling the system to handle both pure quantum algorithms and hybrid quantum-classical programs.

3.6.3 Module 3: Transpilation and Optimization

The Transpilation Module performs hardware-aware circuit compilation, mapping logical quantum circuits to physical hardware topologies with connectivity constraints. This module implements routing algorithms (Greedy BFS and SABRE) that insert SWAP gates to satisfy topology requirements while minimizing circuit depth and maintaining logical equivalence.

Table 3.2: Functional Requirements - Parsing Module

Req. ID	Description
FR-1.1	The system shall parse OpenQASM 2.0 files with support for standard gates (h, x, y, z, cx, measure)
FR-1.2	The system shall parse OpenQASM 3.0 files with support for control flow (if, for, while) and classical operations
FR-1.3	The system shall parse JSON-formatted gate lists with gate names, qubits, and parameters
FR-1.4	The system shall validate qubit indices to ensure they are within declared range
FR-1.5	The system shall validate gate parameters (e.g., rotation angles) for correct data types
FR-1.6	The system shall report syntax errors with line numbers and descriptive messages
FR-1.7	The system shall support classical register declarations (bit, bit[n])
FR-1.8	The system shall support qubit register declarations (qubit, qubit[n])

3.6.4 Module 4: Simulation Engines

The Simulation Module executes quantum circuits using statevector evolution or tensor network methods. This module supports exact simulation for circuits up to 12 qubits using full statevector representation, and approximate simulation for larger low-entanglement circuits using Matrix Product State (MPS) decomposition with configurable bond dimensions.

3.6.5 Module 5: Visualization and Output

The Visualization Module generates graphical representations of quantum circuits and measurement results. This module creates circuit diagrams showing gate operations on qubit wires, probability histograms for measurement outcomes, and exports circuits to OpenQASM 2.0 format for interoperability with external quantum computing platforms.

3.6.6 Module 6: User Interfaces

The User Interface Module provides three access methods for interacting with the QVM system: a command-line interface for scripting and automation, a RESTful web API for programmatic access, and a browser-based graphical interface for interactive circuit composition and execution. These interfaces accommodate different user workflows from batch processing to educational exploration.

Table 3.3: Functional Requirements - IR Module

Req. ID	Description
FR-2.1	The system shall represent quantum circuits as a Quantum-Circuit object with num_qubits and operations list
FR-2.2	The system shall store operations as dictionaries containing gate name, qubits, parameters, and conditions
FR-2.3	The system shall support conditional operations with register, index, and value specifications
FR-2.4	The system shall support measurement operations with target classical bit specification
FR-2.5	The system shall support label and jump operations for control flow
FR-2.6	The system shall support classical operations (assignment, AND, OR, XOR, NOT)
FR-2.7	The system shall support delay operations with duration specifications
FR-2.8	The system shall provide string representation of circuits for debugging

3.7 Non-Functional Requirements

Non-functional requirements define the quality attributes and constraints that the QVM system must satisfy. These requirements ensure that the system not only provides the correct functionality but also meets standards for performance, reliability, usability, portability, maintainability, and security. The following subsections detail the specific non-functional requirements across these six quality dimensions.

3.7.1 Performance Requirements

Performance requirements specify the execution speed and resource efficiency that the system must achieve. These requirements ensure that the QVM can simulate quantum circuits within acceptable time bounds on standard computing hardware, making it practical for educational use and algorithm development.

3.7.2 Reliability Requirements

Reliability requirements define the system's ability to handle errors gracefully and maintain correct operation under various conditions. These requirements ensure that the QVM validates inputs, detects potential issues like infinite loops, and provides clear diagnostic information when problems occur.

3.7.3 Usability Requirements

Usability requirements specify the ease of use and learnability of the system for different user classes. These requirements ensure that the QVM provides clear documentation, intuitive inter-

Table 3.4: Functional Requirements - Transpiler Module

Req. ID	Description
FR-3.1	The system shall support Greedy BFS routing strategy for qubit mapping
FR-3.2	The system shall support SABRE routing strategy with look-ahead heuristics
FR-3.3	The system shall insert SWAP gates to satisfy topology connectivity constraints
FR-3.4	The system shall support linear chain topology (qubits connected in a line)
FR-3.5	The system shall support grid topology (qubits connected in 2D lattice)
FR-3.6	The system shall optionally restore logical-to-physical qubit mapping after transpilation
FR-3.7	The system shall validate that logical circuit fits within target architecture qubit count
FR-3.8	The system shall cache distance calculations for performance optimization
FR-3.9	The system shall decompose complex gates (Toffoli) into native gate sets
FR-3.10	The system shall preserve single-qubit gates without modification during transpilation

faces, and helpful examples that enable users to quickly understand and effectively use the system for quantum algorithm development and education.

3.7.4 Portability Requirements

Portability requirements define the system’s ability to run on different operating systems and integrate with external quantum computing platforms. These requirements ensure that the QVM can be deployed across Windows, Linux, and macOS environments, and can exchange circuit definitions with industry-standard tools through OpenQASM format support.

3.7.5 Maintainability Requirements

Maintainability requirements specify the code quality standards and development practices that facilitate future modifications and extensions. These requirements ensure that the QVM codebase is well-documented, follows consistent coding conventions, and uses modular design principles that enable developers to understand, modify, and extend the system efficiently.

3.7.6 Security Requirements

Security requirements define the system’s protection mechanisms against malicious inputs and resource exhaustion attacks. These requirements ensure that the QVM validates all user inputs,

Table 3.5: Functional Requirements - Simulator Module

Req. ID	Description
FR-4.1	The system shall simulate circuits using exact statevector evolution for up to 12 qubits
FR-4.2	The system shall support single-qubit gates (H, X, Y, Z, Rx, Ry, Rz, S, T, SX)
FR-4.3	The system shall support two-qubit gates (CX, SWAP)
FR-4.4	The system shall support three-qubit gates (CCX/Toffoli)
FR-4.5	The system shall perform measurements with probabilistic state collapse
FR-4.6	The system shall maintain classical memory registers synchronized with quantum state
FR-4.7	The system shall support conditional gate execution based on classical bit values
FR-4.8	The system shall support control flow (labels, jumps, loops) with infinite loop detection
FR-4.9	The system shall support MPS simulation for low-entanglement circuits with configurable bond dimension
FR-4.10	The system shall report truncation error for MPS simulations
FR-4.11	The system shall use NumPy vectorization for efficient matrix operations
FR-4.12	The system shall support seeded random number generation for reproducible measurements

prevents directory traversal vulnerabilities, and implements resource limits to protect against denial-of-service attempts while maintaining usability for legitimate users.

3.8 External Interface Requirements

External interface requirements specify how the QVM system interacts with users, hardware, software components, and external systems. These requirements define the boundaries of the system and the protocols for communication across those boundaries. The following subsections detail the user interfaces, hardware interfaces, software dependencies, and communication protocols that the system must support.

3.8.1 User Interfaces

The system provides three user interfaces:

Command-Line Interface (CLI): A terminal-based interface accepting file paths and command-line flags. Supports batch processing and scripting. Displays text-based output with optional visualization windows.

Web API: A RESTful HTTP API built with FastAPI. Accepts JSON POST requests at `/execute` endpoint. Returns JSON responses with statevector, probabilities, and execution metadata. Supports CORS for cross-origin requests.

Table 3.6: Functional Requirements - Visualization Module

Req. ID	Description
FR-5.1	The system shall generate circuit diagrams using matplotlib with gates on qubit wires
FR-5.2	The system shall display measurement operations with classical bit targets
FR-5.3	The system shall generate probability histograms for measurement outcomes
FR-5.4	The system shall export circuits to OpenQASM 2.0 format
FR-5.5	The system shall save visualizations to PNG files
FR-5.6	The system shall display execution statistics (circuit depth, gate count)

Table 3.7: Functional Requirements - Interface Module

Req. ID	Description
FR-6.1	The system shall provide a command-line interface accepting file paths and options
FR-6.2	The system shall support CLI flags for transpilation (<code>-transpile</code> , <code>-routing</code> , <code>-architecture</code>)
FR-6.3	The system shall support CLI flags for simulation (<code>-simulator</code> , <code>-shots</code> , <code>-seed</code>)
FR-6.4	The system shall provide a REST API with POST <code>/execute</code> endpoint
FR-6.5	The system shall accept JSON circuit definitions via API
FR-6.6	The system shall return JSON responses with statevector, probabilities, and metadata
FR-6.7	The system shall provide a web GUI with circuit editor
FR-6.8	The system shall display real-time execution results in web GUI
FR-6.9	The system shall provide example circuits in web GUI

Web GUI: A browser-based dashboard with circuit editor, execution controls, and result visualization. Built with HTML/CSS/JavaScript. Communicates with backend via REST API. Displays circuit diagrams and probability histograms inline.

3.8.2 Hardware Interfaces

The system runs on standard computing hardware with the following minimum specifications:

- CPU: Dual-core processor (2.0 GHz or higher)
- RAM: 4 GB minimum (8 GB recommended for 12-qubit circuits)
- Storage: 100 MB for installation
- Display: 1024x768 resolution for GUI

Table 3.8: Non-Functional Requirements - Performance

Req. ID	Description
NFR-1.1	The system shall simulate 10-qubit circuits in under 5 seconds on standard hardware
NFR-1.2	The system shall parse OpenQASM 3.0 files in under 1 second for circuits with up to 1000 operations
NFR-1.3	The system shall transpile circuits with under 100 gates in under 2 seconds using SABRE routing
NFR-1.4	The system shall support MPS simulation of 20-qubit low-entanglement circuits
NFR-1.5	The system shall cache distance calculations to avoid redundant BFS searches

Table 3.9: Non-Functional Requirements - Reliability

Req. ID	Description
NFR-2.1	The system shall detect and prevent infinite loops in control flow with configurable operation limits
NFR-2.2	The system shall validate all input circuits before execution
NFR-2.3	The system shall provide descriptive error messages for all failure modes
NFR-2.4	The system shall maintain numerical stability for statevector normalization
NFR-2.5	The system shall achieve 95% test coverage with automated unit tests

3.8.3 Software Interfaces

The system interfaces with the following software components:

Python Runtime: Requires Python 3.10 or higher for execution.

NumPy: Used for linear algebra operations (matrix multiplication, tensor products). Version 1.21+ required.

Matplotlib: Used for visualization (circuit diagrams, histograms). Version 3.5+ required.

Lark: Used for OpenQASM 3.0 parsing with LALR parser. Version 1.1+ required.

FastAPI: Used for web API server. Version 0.95+ required.

Uvicorn: ASGI server for running FastAPI application.

3.8.4 Communication Interfaces

The system uses the following communication protocols:

HTTP/HTTPS: Web API communicates via HTTP POST requests with JSON payloads. Supports standard HTTP status codes (200, 400, 500).

File I/O: Reads circuit files from local filesystem. Supports .qasm, .json file extensions. Writes visualization outputs to .png files.

Standard I/O: CLI uses stdin for input and stdout/stderr for output. Supports piping and

Table 3.10: Non-Functional Requirements - Usability

Req. ID	Description
NFR-3.1	The system shall provide comprehensive documentation with usage examples
NFR-3.2	The system shall provide clear CLI help messages with <code>-help</code> flag
NFR-3.3	The system shall provide example circuits for common algorithms (Bell, GHZ, Grover)
NFR-3.4	The system shall use intuitive gate naming conventions consistent with Qiskit
NFR-3.5	The web GUI shall be accessible without installation via web browser

Table 3.11: Non-Functional Requirements - Portability

Req. ID	Description
NFR-4.1	The system shall run on Windows, Linux, and macOS operating systems
NFR-4.2	The system shall require only Python 3.10+ and standard scientific libraries
NFR-4.3	The system shall export circuits to OpenQASM 2.0 for compatibility with IBM Quantum
NFR-4.4	The system shall support multiple input formats (QASM 2.0, QASM 3.0, JSON)

redirection.

3.9 Requirement Traceability Matrix

3.10 Summary

This chapter has presented a comprehensive analysis of the Quantum Virtual Machine requirements, covering functional requirements across six major modules (Parsing, IR, Transpilation, Simulation, Visualization, Interfaces) and non-functional requirements across six quality attributes (Performance, Reliability, Usability, Portability, Maintainability, Security). The requirements were elicited through literature review, comparative analysis, and iterative prototyping. Ten detailed use cases were documented, covering the primary user interactions with the system. The requirement traceability matrix demonstrates clear alignment between business objectives and functional requirements. These requirements guided the system design and implementation, ensuring that the QVM achieves its goal of providing a hardware-agnostic, educational, and lightweight quantum execution runtime.

Table 3.12: Non-Functional Requirements - Maintainability

Req. ID	Description
NFR-5.1	The system shall follow modular architecture with clear separation of concerns
NFR-5.2	The system shall include docstrings for all public functions and classes
NFR-5.3	The system shall use type hints for function signatures
NFR-5.4	The system shall maintain version control with Git
NFR-5.5	The system shall follow PEP 8 Python style guidelines

Table 3.13: Non-Functional Requirements - Security

Req. ID	Description
NFR-6.1	The system shall validate all file paths to prevent directory traversal attacks
NFR-6.2	The system shall sanitize user input in web API to prevent injection attacks
NFR-6.3	The system shall limit maximum circuit size to prevent resource exhaustion
NFR-6.4	The system shall implement operation count limits to prevent denial-of-service

Table 3.14: Requirement Traceability Matrix

Business Objective	Functional Requirements	Use Cases
BO-1: Hardware Agnosticism	FR-3.1, FR-3.2, FR-3.3, FR-3.4, FR-3.5	UC-1.2
BO-2: Educational Value	FR-5.1, FR-5.2, FR-5.3, FR-6.7, FR-6.9	UC-1.5, UC-1.6, UC-1.10
BO-3: Lightweight Runtime	FR-4.1, FR-4.11, NFR-1.1, NFR-4.2	UC-1.3
BO-4: Interoperability	FR-1.1, FR-1.2, FR-5.4, NFR-4.3	UC-1.1, UC-1.7
BO-5: Flexibility	FR-2.3, FR-2.5, FR-2.6, FR-4.7, FR-4.8	UC-1.3, UC-1.4

Chapter 4

Software Design

4.1 Introduction

This chapter presents the software design of the Quantum Virtual Machine (QVM) system. The design follows object-oriented principles and emphasizes modularity, maintainability, and extensibility. The chapter documents the architectural decisions, design patterns, and structural organization that enable the QVM to achieve its goal of hardware-agnostic quantum circuit execution. The design is presented through multiple perspectives using UML diagrams, including architectural views, behavioral models, and data flow representations.

4.2 Design Methodology

The QVM system was designed using **Object-Oriented Design (OOD)** methodology with emphasis on the following principles:

Separation of Concerns: The system is divided into distinct modules (Parser, IR, Transpiler, Simulator, Visualizer) with well-defined responsibilities. Each module encapsulates a specific aspect of the quantum compilation and execution pipeline.

Pipeline Architecture: The design follows a linear pipeline pattern where data flows through sequential stages of transformation: Input $\rightarrow$ Parsing $\rightarrow$ IR $\rightarrow$ Transpilation $\rightarrow$ Simulation $\rightarrow$ Output. This architecture mirrors the natural workflow of quantum program execution.

Strategy Pattern: The transpiler module implements the Strategy pattern, allowing runtime selection between different routing algorithms (Greedy BFS, SABRE) without modifying client code.

Abstraction and Interfaces: The QuantumCircuit IR serves as an abstraction layer that decouples high-level circuit definitions from low-level execution details, enabling hardware independence.

Modularity and Extensibility: Each component is designed as an independent module with minimal coupling, allowing new parsers, simulators, or transpilation strategies to be added without affecting existing code.

4.3 System Overview

The QVM system architecture is organized as a layered pipeline with clear boundaries between components. Figure 4.1 shows the system context, illustrating how the QVM interacts with external actors (developers, students, administrators) and external systems (file system, web browsers, IBM Quantum platform). The system accepts quantum programs through three interfaces (CLI, Web API, Web GUI) and produces results in multiple formats (statevector, probabilities, visualizations, OpenQASM export).




Figure 4.1: System Context Diagram

The context diagram establishes the system boundary and identifies all external entities that interact with the QVM. Users provide quantum circuits as input and receive execution results as output. The system reads circuit files from the file system and can export results back to files. The web-based interfaces communicate with browsers via HTTP, and the OpenQASM export capability enables integration with external quantum platforms like IBM Quantum.

4.4 Architectural Design

The QVM follows a **7-layer pipeline architecture** as shown in Figure 4.2. Each layer represents a distinct phase in the quantum program lifecycle:

Layer 1 - Input Layer: Accepts quantum programs in three formats (OpenQASM 3.0, OpenQASM 2.0, JSON) through three interfaces (CLI, Web API, Web GUI). This layer provides flexibility in how users interact with the system.

Layer 2 - Parser Layer: Transforms textual circuit descriptions into structured Abstract Syntax Trees (AST). Uses Lark parser generator for OpenQASM 3.0 with formal grammar definition. Validates syntax and semantics before proceeding.

Layer 3 - Intermediate Representation (IR): Converts AST into a hardware-agnostic QuantumCircuit object. This layer decouples the input format from the execution backend, enabling the "Write Once, Run Anywhere" philosophy.

Layer 4 - Decomposer: Breaks down complex gates (Toffoli, controlled gates) into native gate sets supported by target architectures. Maps high-level operations to primitive gates (H, X, Y, Z, CX).

Layer 5 - Transpiler: Maps logical qubits to physical qubits according to hardware topology constraints. Implements two routing strategies (Greedy BFS, SABRE) and inserts SWAP gates to satisfy connectivity requirements.

Layer 6 - Simulator: Executes the transpiled circuit using either exact statevector simulation (up to 12 qubits) or efficient Matrix Product State (MPS) simulation (20+ qubits for low-entanglement circuits). Maintains classical memory for hybrid quantum-classical computation.

Layer 7 - Output Layer: Produces results in multiple formats: final statevector, measurement probabilities, circuit visualizations (matplotlib), and OpenQASM export for external tools.

This layered architecture ensures that each component has a single, well-defined responsibility and can be tested, modified, or replaced independently.

4.5 Activity Diagrams

Activity diagrams model the dynamic behavior of the system by showing the flow of control through various processes. The QVM system has three primary workflows:

4.5.1 Circuit Execution Flow

Figure 4.3 shows the complete workflow from loading a circuit file to displaying results. The process includes error handling for invalid files and syntax errors, backend selection (Statevector vs MPS), optional transpilation, and result visualization.




Figure 4.2: System Architecture Diagram

The execution flow begins with the user providing a circuit file. The system validates the file existence and format, then invokes the appropriate parser. If parsing succeeds, the user can optionally request transpilation for a specific architecture. The system then selects the appropriate simulator based on circuit size and user preference. After simulation, results are displayed and optionally saved to files.

4.5.2 Transpilation Process

Figure 4.4 details the transpilation workflow, showing how logical circuits are mapped to physical hardware topologies. The process includes gate decomposition, qubit routing (with choice between Greedy and SABRE algorithms), SWAP insertion, and optional mapping restoration.




Figure 4.3: Activity Diagram - Circuit Execution Flow

The transpilation process starts by analyzing the logical circuit's connectivity requirements. Complex gates are decomposed into native gates. The routing algorithm then determines the optimal qubit mapping, inserting SWAP gates where necessary to satisfy topology constraints. If mapping restoration is enabled, additional SWAPs are inserted to return qubits to their original logical positions.

4.5.3 Simulation Process

Figure 4.5 illustrates the simulation workflow, including state initialization, gate application, measurement with collapse, and control flow handling (labels, jumps, conditionals).




Figure 4.4: Activity Diagram - Transpilation Process




Figure 4.5: Activity Diagram - Simulation Process

The simulation process initializes the quantum state to $|0\rangle^{\otimes n}$ and classical memory to zeros. It then iterates through circuit operations, applying unitary transformations for gates, performing probabilistic collapse for measurements, and updating classical memory. Control flow operations (if, while, jump) are handled by maintaining a program counter and label map. The process includes infinite loop detection to prevent runaway execution.

4.6 Class Diagram

The class diagram in Figure 4.6 shows the static structure of the QVM system, including all major classes, their attributes, methods, and relationships.




Figure 4.6: Class Diagram - System Structure

Core Classes:

QuantumCircuit: The central IR class that represents a hardware-agnostic quantum circuit. Contains `num_qubits`, `operations` list, and `classical_registers` dictionary. Provides methods for adding operations, classical registers, and string representation.

Simulator: Implements exact statevector simulation using NumPy. Contains gate matrices (H, X, Y, Z, CX, etc.) and methods for applying gates, measurements, and control flow. The `simulate()` method returns the final statevector and classical memory state.

MPSSimulator: Implements Matrix Product State simulation for efficient handling of low-entanglement circuits. Uses tensor contractions and SVD truncation to maintain computational efficiency.

Transpiler: Maps logical circuits to physical architectures. Contains `target_architecture`, `strategy` (greedy/sabre), and `restore_mapping` flag. Implements BFS routing, SABRE heuristics, and SWAP insertion algorithms.

OpenQASM3Parser: Parses OpenQASM 3.0 files using Lark grammar. Implements two-pass parsing: first pass for declarations, second pass for operations. Handles control flow (if, for, while) and classical operations.

QASMParser: Parses OpenQASM 2.0 files with basic gate support. Simpler than OpenQASM3Parser but sufficient for standard quantum algorithms.

TargetArchitecture: Represents hardware topology with `num_qubits` and `connectivity` list. Provides `is_connected()` method to check if two qubits are adjacent.

Decomposer: Breaks down complex gates into native gate sets. Implements decomposition rules (e.g., Toffoli $\rightarrow$ CNOTs + single-qubit gates).

Visualizer: Generates circuit diagrams and probability histograms using matplotlib. Provides `draw_circuit()` and `plot_probabilities()` methods.

CLI: Command-line interface that parses arguments and orchestrates the execution pipeline. Handles file loading, transpilation options, and result display.

APIServer: FastAPI-based REST API server. Provides /execute endpoint that accepts JSON circuit definitions and returns JSON results.

Relationships: - Simulator and MPSSimulator both depend on QuantumCircuit - Transpiler depends on QuantumCircuit and TargetArchitecture - Parsers create QuantumCircuit objects - CLI and APIServer orchestrate all components

4.7 Sequence Diagrams

Sequence diagrams show the dynamic interactions between objects over time. The QVM system has three key interaction scenarios:

4.7.1 CLI Execution Flow

Figure 4.7 shows the interaction sequence when a user executes a circuit via the command-line interface.

The sequence begins with the user invoking the CLI with a circuit file. The CLI loads the file and delegates to the Parser, which creates a QuantumCircuit IR. If transpilation is requested, the CLI invokes the Transpiler with the target architecture. The Transpiler returns a physical circuit with SWAP gates inserted. The CLI then invokes the Simulator, which executes the circuit and returns the statevector and classical memory. Finally, the CLI invokes the Visualizer to display results to the user.

4.7.2 API Request Flow

Figure 4.8 illustrates the interaction when a client submits a circuit via the REST API.




Figure 4.7: Sequence Diagram - CLI Execution

The client sends an HTTP POST request with a JSON circuit definition. The APIServer validates the request format and extracts the circuit data. It then delegates to the Parser to create a QuantumCircuit IR. The APIServer invokes the Simulator to execute the circuit. After receiving the results, the APIServer formats them as JSON and returns an HTTP 200 response to the client. Error handling is included for invalid JSON (HTTP 400) and execution errors (HTTP 500).

4.7.3 Transpilation Detail

Figure 4.9 shows the detailed interaction during the transpilation process.




Figure 4.8: Sequence Diagram - API Request

The Transpiler receives a logical circuit and target architecture. It first invokes the Decomposer to break down complex gates. The Decomposer returns a circuit with only native gates. The Transpiler then analyzes connectivity requirements and selects the routing strategy (Greedy or SABRE). For each two-qubit gate that violates topology constraints, the Transpiler inserts SWAP gates to bring qubits into adjacency. The process continues until all gates satisfy connectivity requirements. If mapping restoration is enabled, additional SWAPs are inserted to restore the original logical-to-physical mapping.

4.8 State Transition Diagrams

State transition diagrams model the lifecycle of key entities in the system. Figure 4.10 shows the state transitions for a quantum circuit as it progresses through the QVM pipeline.




Figure 4.9: Sequence Diagram - Transpilation Detail

Circuit States:

Created: Initial state after parsing. The circuit exists as a QuantumCircuit IR object but has not been validated or processed.

Validated: The circuit has passed semantic validation (qubit indices in range, valid gate parameters, classical registers declared).

Decomposed: Complex gates have been broken down into native gate sets. The circuit now contains only primitive operations.

Transpiled: The circuit has been mapped to a physical architecture with SWAP gates inserted. Logical qubits are mapped to physical qubits.

Executing: The simulator is actively processing the circuit operations. The quantum state is being evolved step-by-step.

Completed: Simulation has finished successfully. The final statevector and classical memory are available.

Error: An error occurred during any stage (parsing, validation, transpilation, or execution). The circuit cannot proceed further.

State Transitions: - Created → Validated: Semantic validation succeeds - Created → Error: Validation fails (invalid qubit index, undefined register) - Validated → Decomposed: Gate decomposition succeeds - Decomposed → Transpiled: Transpilation succeeds (or skipped if not requested) - Transpiled → Executing: Simulation begins - Executing → Completed: All operations executed successfully - Executing → Error: Runtime error (infinite loop, invalid operation)

4.9 Data Flow Diagrams

4.9.1 Level 1 DFD - System Level

Figure 4.11 shows the high-level data flows between the QVM system and external entities.

Circuit Lifecycle State Transition Diagram




Figure 4.10: State Transition Diagram - Circuit Lifecycle




Figure 4.11: Data Flow Diagram - Level 1 (System Level)

External Entities: - User: Provides circuit files and execution parameters - File System: Stores circuit files and output results - Web Browser: Displays web GUI and receives results

Data Flows: - Circuit File → QVM System: User provides quantum program - Execution Parameters → QVM System: User specifies options (architecture, simulator, seed) - Results → User: System returns statevector, probabilities, visualizations - Circuit File ← File System: System reads input files - Output Files → File System: System writes visualizations and exports - HTTP Request → QVM System: Browser sends circuit via API - HTTP Response → Web Browser: System returns JSON results

Data Stores: - Circuit Files: Persistent storage of quantum programs - Configuration: System settings and architecture definitions

4.9.2 Level 2 DFD - Module Level

Figure 4.12 shows the detailed data flows between internal modules.




Figure 4.12: Data Flow Diagram - Level 2 (Module Level)

Processes:

P1 - Parse Circuit: Accepts circuit text and produces QuantumCircuit IR. Validates syntax and semantics.

P2 - Decompose Gates: Accepts QuantumCircuit with complex gates and produces QuantumCircuit with native gates only.

P3 - Transpile Circuit: Accepts logical QuantumCircuit and target architecture, produces physical QuantumCircuit with SWAP gates.

P4 - Simulate Circuit: Accepts QuantumCircuit and execution parameters, produces statevector and classical memory.

P5 - Visualize Results: Accepts statevector and circuit, produces diagrams and histograms.

P6 - Export Circuit: Accepts QuantumCircuit, produces OpenQASM string.

Data Flows: - Circuit Text → P1: Raw input from file or API - QuantumCircuit IR → P2: Parsed circuit with all gates - Native Gate Circuit → P3: Circuit with only primitive gates - Physical Circuit → P4: Transpiled circuit ready for execution - Statevector + Classical Memory → P5: Simulation results - QuantumCircuit IR → P6: Circuit to be exported - OpenQASM String → User: Exported circuit for external tools

Data Stores: - Architecture Definitions: Topology specifications (linear, grid) - Gate Decomposition Rules: Mapping from complex gates to native gates - Simulation State: Intermediate quantum state during execution

4.10 Entity-Relationship Diagram

The QVM system does not use a traditional relational database. Instead, it uses a **file-based data model** where quantum circuits are stored as text files (OpenQASM, JSON) and results are stored as output files (PNG images, QASM exports).

Rationale for File-Based Approach:

- **Simplicity:** Quantum circuits are naturally represented as text files, making them human-readable and version-controllable.

- **Portability:** File-based storage allows easy sharing of circuits between users and systems without database dependencies.
- **Stateless Execution:** Each circuit execution is independent, requiring no persistent state between runs.
- **Educational Focus:** File-based approach is easier for students to understand and debug compared to database abstractions.

File Structure:

- **Input Files:** .qasm (OpenQASM 2.0/3.0), .json (JSON gate lists)
- **Output Files:** .png (visualizations), .qasm (exports), .txt (results)
- **Configuration Files:** .json (architecture definitions)

If the system were to be extended with a database for storing execution history or user accounts, a simple schema would include:

- **Users:** user_id, username, email
- **Circuits:** circuit_id, user_id, name, qasm_text, created_at
- **Executions:** execution_id, circuit_id, architecture, simulator, results, timestamp

However, this is beyond the current scope of the educational QVM system.

4.11 Data Dictionary

The data dictionary documents the structure of key data elements in the QVM system.

4.11.1 QuantumCircuit IR Structure

This section defines the internal representation for the quantum circuit, storing the number of qubits and sequences of operations.

Table 4.1: QuantumCircuit Data Structure

Attribute	Data Type	Description
num_qubits	int	Number of qubits in the circuit (must be positive integer)
operations	list[dict]	Ordered list of quantum and classical operations
classical_registers	dict[str, int]	Mapping from register name to size

4.11.2 Operation Dictionary Structure

The operation dictionary specifies each individual quantum or classical operation within a circuit.

Table 4.2: Operation Data Structure

Field	Data Type	Description
name	str	Gate name (h, x, cx, measure, etc.)
qubits	list[int]	Target qubit indices (0-indexed)
params	list[float]	Gate parameters (e.g., rotation angles)
condition	dict — None	Conditional execution specification
target_bit	tuple — None	(register_name, index) for measurement results
duration	str — None	Timing specification (e.g., "100ns")
label	str — None	Label name for control flow
jump_to	str — None	Target label for jump operations
classical_op	dict — None	Classical operation specification

4.11.3 Condition Dictionary Structure

The condition structure defines the conditions under which an operation should be conditionally executed based on classical bits.

Table 4.3: Condition Data Structure

Field	Data Type	Description
register	str	Classical register name
index	int	Bit index within register
value	int	Expected value (0 or 1) for condition to be true

4.11.4 Classical Operation Dictionary Structure

Classical operations define operations performed on classical registers such as logical AND, OR, and XOR.

Table 4.4: Classical Operation Data Structure

Field	Data Type	Description
op	str	Operator (=, &, —, ^, ~)
target	tuple	(register_name, index) for result storage
args	list	Operands (tuples for register references, ints for literals)

4.11.5 Target Architecture Structure

This structure represents the physical architecture of the hardware, including its qubit connectivity constraints.

Table 4.5: TargetArchitecture Data Structure

Attribute	Data Type	Description
num_qubits	int	Total number of physical qubits
connectivity	list[tuple]	List of (qubit1, qubit2) pairs representing edges
name	str	Architecture name (e.g., "linear_5", "grid_3x3")

4.11.6 Simulation Results Structure

The simulation results capture the final execution outcomes, including the statevector and measurement probabilities.

Table 4.6: Simulation Results Data Structure

Field	Data Type	Description
statevector	ndarray	Complex-valued array of size 2^n representing quantum state
classical_memory	dict[str, ndarray]	Mapping from register name to integer array of bit values
probabilities	dict[str, float]	Mapping from basis state to measurement probability
execution_time	float	Time taken for simulation (seconds)

4.12 Summary

This chapter has presented a comprehensive software design for the Quantum Virtual Machine system. The design follows object-oriented principles with a clear 7-layer pipeline architecture that separates concerns and enables modularity. The system structure was documented through multiple UML perspectives: context diagram (system boundary), architecture diagram (layered pipeline), class diagram (static structure), activity diagrams (behavioral workflows), sequence diagrams (object interactions), state diagram (circuit lifecycle), and data flow diagrams (information flow). The design emphasizes hardware abstraction through the QuantumCircuit IR, flexibility through the Strategy pattern for transpilation, and extensibility through modular components. The file-based data model provides simplicity and portability suitable for an educational quantum computing platform. The data dictionary documents the structure of all key data elements, ensuring clear understanding of the system's information architecture. This design provides a solid foundation for the implementation presented in the next chapter.

Chapter 5

Implementation

5.1 Development Environment

The Quantum Virtual Machine was developed using a modern Python-based development environment optimized for scientific computing and quantum algorithm development.

Programming Language: Python 3.10+ was selected as the primary implementation language due to its extensive scientific computing ecosystem, readability, and widespread adoption in the quantum computing community. Python’s dynamic typing and interactive development capabilities accelerated prototyping and testing.

Integrated Development Environment: Visual Studio Code (VS Code) served as the primary IDE, providing intelligent code completion, integrated debugging, Git version control, and terminal access. The Python extension for VS Code enabled linting, type checking, and test discovery.

Version Control: Git was used for version control with GitHub as the remote repository. The project followed a feature-branch workflow with descriptive commit messages. All code changes were tracked, enabling rollback and collaboration.

Package Management: pip was used for dependency management with a requirements.txt file specifying exact versions of all dependencies. This ensures reproducible builds across different development environments.

Testing Framework: pytest was used for automated testing with 36 unit tests covering all major components. Tests were organized by module and executed automatically before commits.

Documentation Tools: Docstrings followed the Google Python Style Guide. Type hints were used throughout for improved code clarity and IDE support. Markdown was used for technical documentation.

Development Workflow:

1. Write code with type hints and docstrings
2. Run unit tests locally (pytest)
3. Commit changes with descriptive messages
4. Push to GitHub remote repository
5. Review code and documentation

5.2 Core Module Implementation

5.2.1 Module 1: Parser Layer

The parser layer is responsible for converting textual quantum circuit descriptions into the internal QuantumCircuit IR. Three parsers were implemented to support different input formats.

OpenQASM 2.0 Parser (src/qvm/parser.py): Implements a basic parser for OpenQASM 2.0 files using string parsing and regular expressions. Supports standard gates (h, x, y, z, cx, measure) and qubit/classical register declarations. The parser validates syntax and creates QuantumCircuit objects with operations list.

OpenQASM 3.0 Parser (src/qvm/qasm3_parser.py): Implements a full AST-based parser using the Lark parser generator. The grammar is defined in qasm3.lark using EBNF notation. The parser performs two passes: first pass identifies declarations (qubits, classical registers), second pass processes operations. Supports advanced features including control flow (if, for, while), classical operations (AND, OR, XOR, NOT), timing (delay), and labels/jumps.

JSON Parser (src/qvm/parser.py): Implements a simple parser for JSON-formatted gate lists. Each gate is represented as a dictionary with name, qubits, and optional parameters. This format is useful for programmatic circuit generation and testing.

Key Implementation Details:

- Two-pass parsing for OpenQASM 3.0 ensures all declarations are processed before operations
- Recursive AST traversal handles nested control structures
- Error messages include line numbers and descriptive text
- Qubit mapping translates named registers to physical indices

5.2.2 Module 2: Intermediate Representation

The IR module (src/qvm/ir.py) defines the QuantumCircuit class, which serves as the hardware-agnostic representation of quantum programs.

QuantumCircuit Class:

- **Attributes:** num_qubits (int), operations (list[dict]), classical_registers (dict[str, int])
- **Methods:** add_classical_register(), add_operation(), __str__()
- **Operation Structure:** Each operation is a dictionary containing name, qubits, params, condition, target_bit, duration, label, jump_to, classical_op

Design Decisions:

- Dictionary-based operations provide flexibility for different gate types
- Classical registers stored separately from quantum state
- Validation performed in add_operation() to catch errors early
- String representation enables debugging and logging

Classical Memory Integration: The IR was extended to support classical registers and conditional operations, enabling hybrid quantum-classical computation. Classical registers are declared with name and size. Operations can include conditions that check classical bit values before execution. Measurement operations specify target classical bits for storing results.

5.2.3 Module 3: Transpiler

The transpiler module (`src/qvm/transpiler.py`) maps logical circuits to physical hardware topologies by inserting SWAP gates to satisfy connectivity constraints.

Transpiler Class:

- **Attributes:** `target_architecture`, `strategy` (`greedy/sabre`), `restore_mapping` (`bool`), `distance_cache` (`dict`)
- **Methods:** `transpile()`, `_transpile_greedy()`, `_transpile_sabre()`, `_swap_update_maps()`, `_bfs_shortest_path`

Greedy BFS Routing: The greedy strategy uses breadth-first search to find the shortest path between disconnected qubits. For each two-qubit gate that violates topology constraints, the algorithm:

1. Finds shortest path using BFS
2. Inserts SWAP gates along the path
3. Updates qubit mapping
4. Executes the gate
5. Optionally swaps back to restore original mapping

SABRE Routing: The SABRE strategy uses a lookahead heuristic to minimize total SWAP count. The algorithm maintains a ready queue of two-qubit gates and evaluates the cost of each potential SWAP using a decay-weighted distance metric. Single-qubit gates are emitted immediately without routing.

Key Implementation Details:

- Distance caching avoids redundant BFS searches
- Qubit mapping maintained with forward (`logical`→`physical`) and inverse (`physical`→`logical`) dictionaries
- Strategy pattern allows runtime selection of routing algorithm
- Optional mapping restoration trades circuit depth for logical consistency

5.2.4 Module 4: Simulator Engines

Two simulation engines were implemented to handle different circuit scales and entanglement profiles.

Statevector Simulator (`src/qvm/simulator.py`): Implements exact quantum simulation using full statevector representation. The state is stored as a complex-valued NumPy array of size 2^n . Gates are applied using matrix-vector multiplication with Kronecker products for multi-qubit operations.

Gate Application Methods:

- **Single-qubit gates:** Construct full operator using Kronecker products, apply via matrix-vector multiplication
- **CX gate:** Use index permutation for efficient application without full matrix construction
- **SWAP gate:** Use index permutation to exchange qubit states
- **CCX (Toffoli) gate:** Use index masking to flip target when both controls are 1

Measurement with Collapse: Measurements are implemented using Born rule probabilities. The algorithm:

1. Compute probability of measuring 0: $P(0) = \sum_{i:q_i=0} |\psi_i|^2$
2. Sample outcome using random number generator
3. Collapse statevector by zeroing amplitudes inconsistent with outcome
4. Renormalize statevector to maintain unit norm
5. Store result in classical memory

Control Flow Handling: The simulator maintains a program counter (PC) and label map to support jumps and loops. Labels are pre-scanned before execution. Jump operations conditionally update the PC based on classical bit values. An operation counter prevents infinite loops by raising an error after `max_ops` iterations.

MPS Simulator (`src/qvm/mps_simulator.py`): Implements Matrix Product State simulation for low-entanglement circuits. The quantum state is represented as a chain of tensors with bond dimension χ . Single-qubit gates are applied via tensor contraction. Two-qubit gates require:

1. Contract adjacent tensors into single tensor
2. Apply two-qubit gate matrix
3. Perform Singular Value Decomposition (SVD)
4. Truncate singular values to maintain bond dimension
5. Split back into two tensors

Truncation Strategy: The bond dimension χ controls the tradeoff between accuracy and efficiency. Small χ (16-32) enables simulation of 20+ qubits but introduces truncation error. Large χ approaches exact simulation but loses efficiency. The simulator tracks cumulative truncation error to inform users of approximation quality.

5.2.5 Module 5: Decomposer

The decomposer module (`src/qvm/decomposer.py`) breaks down complex gates into native gate sets supported by target architectures.

Decomposer Class:

- **Attributes:** `native_gates` (set of supported gate names)
- **Methods:** `decompose_circuit()`, `decompose_toffoli()`, `decompose_controlled_gate()`

Toffoli Decomposition: The Toffoli (CCX) gate is decomposed into 6 CX gates and 9 single-qubit gates using the standard decomposition from Nielsen & Chuang. This enables execution on hardware that only supports CX as the native two-qubit gate.

Controlled Gate Decomposition: Arbitrary controlled gates are decomposed using the general controlled-U decomposition, which expresses a controlled-U as a sequence of CX gates and single-qubit rotations.

5.2.6 Module 6: Visualization

The visualization module (`src/qvm/visual.py`) generates circuit diagrams and probability histograms using matplotlib.

Circuit Diagram Generation: Circuits are rendered as horizontal qubit wires with gates drawn as boxes or symbols. The algorithm:

1. Draw horizontal lines for each qubit
2. Iterate through operations in order
3. Draw single-qubit gates as boxes on appropriate wire
4. Draw two-qubit gates as boxes with vertical connecting lines
5. Draw measurement operations as meter symbols
6. Label gates with names and parameters

Probability Histogram: Measurement probabilities are displayed as bar charts with basis states on x-axis and probabilities on y-axis. Only states with probability $> 10^{-6}$ are shown to reduce clutter.

5.3 Algorithm Implementation

5.3.1 SABRE Routing Algorithm

The SABRE (SWAP-based Bidirectional heuristic search) algorithm minimizes SWAP gate count using lookahead heuristics.

Pseudocode:

Algorithm: SABRE Routing

Input: logical_circuit, target_architecture

Output: physical_circuit with SWAP gates

1. Initialize qubit_map = {i: i for all logical qubits}
2. Initialize ready_queue = []
3. Initialize op_index = 0

4. Function enqueue_ready(start_idx):
 For j from start_idx to end of operations:
 If operation[j] is two-qubit gate:
 Add (j, operation[j]) to ready_queue
 Return j + 1
 Else:
 Emit operation[j] to physical_circuit
 Return length(operations)

5. op_index = enqueue_ready(0)

6. While ready_queue is not empty:
 op = ready_queue[0]
 q1, q2 = op.qubits (logical)
 p1, p2 = qubit_map[q1], qubit_map[q2]

 If architecture.is_connected(p1, p2):
 Emit op to physical_circuit
 Remove op from ready_queue
 op_index = enqueue_ready(op_index)
 Continue

best_swap = None

best_cost = infinity

For each edge (s1, s2) in architecture.connectivity:
 Temporarily apply swap(s1, s2) to qubit_map
 cost = heuristic_cost(ready_queue, qubit_map)
 If cost < best_cost:
 best_cost = cost
 best_swap = (s1, s2)
 Revert swap

Emit swap(best_swap) to physical_circuit
Apply best_swap to qubit_map

7. If restore_mapping:

```

For each logical qubit i:
    While qubit_map[i] != i:
        Find swap to move qubit closer to target
        Emit swap to physical_circuit
        Update qubit_map

```

8. Return physical_circuit

Heuristic Cost Function: The cost of a qubit mapping is computed as:

$$\text{cost} = \sum_{i=0}^{\min(3, |\text{ready}|)} \text{decay}^i \cdot \text{distance}(p_1^{(i)}, p_2^{(i)})$$

where $p_1^{(i)}, p_2^{(i)}$ are the physical qubits for the i -th gate in the ready queue, and $\text{decay} = 0.6$ weights near-term gates more heavily.

5.3.2 BFS Shortest Path Algorithm

Breadth-first search finds the shortest path between two qubits in the connectivity graph.

Pseudocode:

Algorithm: BFS Shortest Path

Input: start_node, end_node, connectivity_graph

Output: shortest_path (list of nodes)

1. If start_node == end_node:
 Return [start_node]
2. Build adjacency list from connectivity_graph
3. Initialize queue = [(start_node, [start_node])]
4. Initialize visited = {start_node}
5. While queue is not empty:
 node, path = queue.pop_front()

 For each neighbor in adjacency[node]:
 If neighbor == end_node:
 Return path + [neighbor]

 If neighbor not in visited:
 Add neighbor to visited
 Add (neighbor, path + [neighbor]) to queue
6. Return [] (no path found)

Time Complexity: $O(V + E)$ where V is the number of qubits and E is the number of edges in the connectivity graph.

5.3.3 MPS Tensor Contraction Algorithm

Matrix Product State simulation applies gates via tensor contractions and SVD truncation.

Pseudocode for Two-Qubit Gate:

Algorithm: Apply Two-Qubit Gate (MPS)

Input: tensors (list of MPS tensors), gate_matrix, qubit1, qubit2

Output: updated tensors

1. If $|qubit1 - qubit2| > 1$:
Raise error (non-adjacent qubits require SWAPs)
2. Let $i = \min(qubit1, qubit2)$
3. Contract tensors[i] and tensors[i+1]:
theta = contract(tensors[i], tensors[i+1])
Shape: (bond_left, 2, 2, bond_right)
4. Reshape theta to matrix:
theta_matrix = reshape(theta, (bond_left * 2, 2 * bond_right))
5. Apply gate:
theta_matrix = gate_matrix @ theta_matrix
6. Perform SVD:
U, S, Vt = SVD(theta_matrix)
7. Truncate to bond dimension chi:
U = U[:, :chi]
S = S[:chi]
Vt = Vt[:chi, :]
truncation_error = sum(S[chi:])
8. Reconstruct tensors:
tensors[i] = reshape(U, (bond_left, 2, chi))
tensors[i+1] = reshape(S @ Vt, (chi, 2, bond_right))
9. Return tensors, truncation_error

Truncation Error: The cumulative truncation error is tracked as the sum of discarded singular values across all gate applications.

5.4 External APIs and SDKs

NumPy Integration: NumPy is the foundation of the QVM's computational engine. All quantum states are represented as NumPy arrays with dtype=complex128. Gate matrices are

Table 5.1: External Libraries and Their Usage

Library	Version	Usage in QVM
NumPy	1.21+	Linear algebra operations (matrix multiplication, Kronecker products, SVD), statevector storage, tensor operations
Matplotlib	3.5+	Circuit diagram rendering, probability histogram plotting, visualization of results
Lark	1.1+	OpenQASM 3.0 parsing with LALR parser, AST generation, grammar-based syntax validation
FastAPI	0.95+	REST API server implementation, request validation, JSON serialization, async endpoint handling
Uvicorn	0.20+	ASGI server for running FastAPI application, handles HTTP requests, WebSocket support
Pytest	7.0+	Unit testing framework, test discovery, fixtures, assertion introspection

pre-computed as NumPy arrays and cached for efficiency. Vectorized operations (element-wise multiplication, matrix-vector products) leverage NumPy’s optimized C implementations for performance.

Lark Parser Generator: Lark provides a declarative approach to parsing using EBNF grammar definitions. The QVM’s OpenQASM 3.0 grammar (`qasm3.lark`) defines the syntax rules. Lark generates an AST that the parser traverses to build the IR. The LALR parser provides good performance for the OpenQASM 3.0 grammar.

FastAPI Framework: FastAPI provides automatic request validation using Pydantic models, automatic OpenAPI documentation generation, and `async/await` support for concurrent request handling. The QVM API defines a single `/execute` endpoint that accepts JSON circuit definitions and returns simulation results.

5.5 User Interface Implementation

5.5.1 Command-Line Interface (CLI)

The CLI (`src/qvm/cli.py`) provides a terminal-based interface for executing quantum circuits.

Argument Parsing: The CLI uses `argparse` to define command-line options:

- **input_file:** Path to circuit file (QASM or JSON)
- **-nqubits:** Number of qubits (required for JSON)
- **-transpile:** Enable transpilation
- **-routing:** Routing strategy (greedy/sabre)

- **–no-restore-mapping:** Disable mapping restoration
- **–visualize:** Show plots
- **–seed:** RNG seed for reproducibility

Execution Flow:

1. Parse command-line arguments
2. Load circuit file (detect format from extension)
3. Parse circuit into IR
4. Decompose complex gates
5. Optionally transpile for target architecture
6. Simulate circuit
7. Display results (probabilities, classical memory)
8. Optionally show visualizations

Error Handling: The CLI catches exceptions at each stage and displays user-friendly error messages with context. File not found errors, parsing errors, and simulation errors are all handled gracefully with `sys.exit(1)`.

5.5.2 Web API

The Web API (`api/app.py`) provides a REST interface for programmatic access.

Endpoint: **POST** `/execute`

- **Request Body:** JSON object with circuit definition, `num_qubits`, `options`
- **Response:** JSON object with `statevector`, `probabilities`, `classical_memory`, `execution_time`
- **Status Codes:** 200 (success), 400 (invalid request), 500 (execution error)

Request Validation: FastAPI automatically validates request bodies using Pydantic models. Invalid JSON or missing required fields result in automatic 422 responses with detailed error messages.

CORS Support: The API includes CORS middleware to allow cross-origin requests from web browsers, enabling the web GUI to communicate with the API.

5.5.3 Web GUI

The Web GUI (`web/index.html`) provides an interactive browser-based interface.

Features:

- Circuit editor with syntax highlighting
- Example circuits (Bell state, GHZ, Grover)
- Execution options (transpilation, routing strategy)
- Real-time result display (circuit diagram, histogram)
- Error message display

Implementation: The GUI is implemented as a single-page application using vanilla JavaScript. Circuit text is sent to the API via `fetch()` POST request. Results are rendered using HTML canvas for circuit diagrams and Chart.js for probability histograms.

Screenshots:

Figure 5.1 shows the Web GUI main interface with the circuit editor, example circuits drop-down, and execution options. The interface provides a clean, intuitive layout for quantum circuit composition and execution.

Figure 5.2 displays the execution results including the circuit visualization and probability histogram. The circuit diagram shows the quantum gates applied to each qubit, while the histogram visualizes the measurement outcome distribution.

5.5.4 Command-Line Interface

The CLI provides direct access to QVM functionality through terminal commands, suitable for scripting, automation, and integration with other tools.

Features:

- File-based circuit loading (QASM, JSON)
- Configurable transpilation options
- Multiple output formats (text, JSON, visualization)
- Verbose mode for debugging
- Batch processing support

Figure 5.3 demonstrates CLI execution with multiple quantum algorithms including Bell state preparation, GHZ state generation, and other quantum circuits. The output shows circuit details, execution results, measurement probabilities, and classical memory states.

5.6 Deployment

5.6.1 Local Installation

The QVM is designed for local installation and execution without cloud dependencies.

Installation Steps:

1. Clone repository: `git clone https://github.com/qayum/quantum-virtual-machine.git`
2. Navigate to directory: `cd quantum-virtual-machine`
3. Install dependencies: `pip install -r requirements.txt`
4. Verify installation: `python -m pytest tests/`

Requirements:

- Python 3.10 or higher
- 4 GB RAM minimum (8 GB recommended for 12-qubit circuits)
- 100 MB disk space
- Operating System: Windows, Linux, or macOS

5.6.2 Running the System

CLI Usage:

Run a QASM file

```
python -m src.qvm.cli examples/bell_state.qasm --visualize
```

Run with transpilation

```
python -m src.qvm.cli examples/grover_101.json --nqubits 3 \
--transpile --routing sabre --visualize
```

API Server:

Start server

```
python -m src.qvm.server --host 127.0.0.1 --port 8000
```

Access GUI

Open browser to `http://127.0.0.1:8000`

5.6.3 Testing

The QVM includes a comprehensive test suite with 36 unit tests covering all major components.

Running Tests:

```
# Run all tests
python -m pytest tests/

# Run specific test file
python -m pytest tests/test_simulator.py


# Run with coverage report
python -m pytest --cov=src/qvm tests/
```

Test Categories:

- Parser tests: Validate QASM parsing and IR generation
- Simulator tests: Verify gate application and measurement
- Transpiler tests: Check SWAP insertion and routing
- Integration tests: End-to-end circuit execution
- Algorithm tests: Verify Bell, GHZ, Bernstein-Vazirani, Grover

5.7 Summary

This chapter has documented the implementation of the Quantum Virtual Machine system. The development environment utilized Python 3.10+ with modern tools (VS Code, Git, pytest) for efficient development and testing. Six core modules were implemented: Parser (OpenQASM 2.0/3.0/JSON), Intermediate Representation (QuantumCircuit class), Transpiler (Greedy and SABRE routing), Simulator (Statevector and MPS engines), Decomposer (gate breakdown), and Visualization (matplotlib-based rendering). Three key algorithms were detailed with pseudocode: SABRE routing (lookahead heuristic), BFS shortest path (connectivity analysis), and MPS tensor contraction (SVD truncation). External libraries (NumPy, Matplotlib, Lark, FastAPI) were integrated for computational efficiency and API functionality. Three user interfaces were implemented: CLI (argparse-based), Web API (FastAPI REST endpoint), and Web GUI (browser-based SPA). The system supports local deployment with simple installation via pip and includes 36 automated tests for quality assurance. The implementation successfully realizes the design presented in Chapter 4, providing a complete hardware-agnostic quantum execution runtime.




../Screenshots/RuningExapmle_WebUI01.png

Figure 5.1: Web GUI Main Interface - Circuit Editor and Execution Options



../Screenshots/RuningExapmle_WebUI02.png

Figure 5.2: Web GUI Results Display - Circuit Visualization and Probability Histogram

The image is a large rectangular box that is mostly empty, representing a screenshot of a CLI execution. The only visible text within the box is the file path `../Screenshots/CLI_Output.png` located on the left side.

`../Screenshots/CLI_Output.png`

Figure 5.3: CLI Execution Examples - Multiple Quantum Algorithm Executions

Chapter 6

Testing and Evaluation

6.1 Testing Strategy

The Quantum Virtual Machine employs a comprehensive testing strategy to ensure correctness, reliability, and performance. The testing approach follows the standard software testing pyramid with unit tests forming the foundation, integration tests validating component interactions, and system tests verifying end-to-end functionality.

Testing Framework: pytest was selected as the primary testing framework due to its simplicity, powerful fixtures, and excellent assertion introspection. All tests are located in the tests/ directory and can be executed with a single command.

Test Organization: Tests are organized by module with separate test files for each major component:

- test_parser.py - Parser validation
- test_simulator.py - Simulator correctness
- test_transpiler.py - Transpilation logic
- test_decomposer.py - Gate decomposition
- test_ir.py - Intermediate representation
- test_api.py - REST API endpoints
- test_visual.py - Visualization functions

Test Execution: Tests are executed automatically before commits and can be run with coverage reporting:

```
# Run all tests
python -m pytest tests/
```

```
# Run with coverage
python -m pytest --cov=src/qvm tests/
```

Continuous Testing: The project follows a test-driven development approach where tests are written before or alongside implementation code. This ensures that all features have corresponding test coverage and that regressions are caught early.

6.2 Unit Testing

Unit tests validate individual functions and methods in isolation. The QVM includes 36 unit tests covering all major components.

6.2.1 Parser Tests

Table 6.1: Parser Unit Tests

Test ID	Test Description	Status	Expected Result
UT-P-01	Parse valid circuit with H, CX, RZ gates	Pass	QuantumCircuit with 3 operations
UT-P-02	Parse empty circuit description	Pass	QuantumCircuit with 0 operations
UT-P-03	Parse circuit missing 'name' field	Pass	ValueError raised
UT-P-04	Parse circuit missing 'qubits' field	Pass	ValueError raised
UT-P-05	Parse circuit with invalid qubit index	Pass	ValueError raised

Test Coverage: Parser tests achieve 95% code coverage, validating both success and error paths. Edge cases include empty circuits, missing fields, and invalid qubit indices.

6.2.2 Simulator Tests

Table 6.2: Simulator Unit Tests

Test ID	Test Description	Status	Expected Result
UT-S-01	Apply Hadamard gate to $ 0\rangle$	Pass	$\frac{1}{\sqrt{2}}(0\rangle + 1\rangle)$
UT-S-02	Apply Pauli-X gate to $ 0\rangle$	Pass	$ 1\rangle$
UT-S-03	Apply $\text{RY}(\pi/2)$ gate to $ 0\rangle$	Pass	$\frac{1}{\sqrt{2}}(0\rangle + 1\rangle)$
UT-S-04	Create Bell state (H + CX)	Pass	$\frac{1}{\sqrt{2}}(00\rangle + 11\rangle)$
UT-S-05	Create GHZ state (H + 2 CX)	Pass	$\frac{1}{\sqrt{2}}(000\rangle + 111\rangle)$
UT-S-06	Handle unsupported gate	Pass	ValueError raised
UT-S-07	Handle invalid qubit count for gate	Pass	ValueError raised
UT-S-08	Sample Bell state (2000 shots)	Pass	$\approx 50\% \text{ } 00\rangle, \approx 50\% \text{ } 11\rangle$
UT-S-09	Measurement collapse to basis state	Pass	Single basis state, norm=1

Validation Method: Simulator tests compare computed statevectors against analytically known results using `np.allclose()` with tolerance 10^{-7} . Probability distributions are validated using statistical sampling with 2000 shots and tolerance $\pm 5\%$.

Key Test Cases:

- **Single-qubit gates:** Verify H, X, Y, Z, RY gates produce correct statevectors
- **Two-qubit gates:** Verify CX gate creates entanglement correctly
- **Multi-qubit circuits:** Verify Bell and GHZ states have correct probability distributions
- **Measurement:** Verify measurement collapses statevector and maintains normalization
- **Error handling:** Verify unsupported gates and invalid qubit counts raise appropriate errors

6.2.3 Transpiler Tests

Table 6.3: Transpiler Unit Tests

Test ID	Test Description	Status	Expected Result
UT-T-01	Transpile on fully connected architecture	Pass	No SWAP gates inserted
UT-T-02	Transpile CX(0,2) on linear chain	Pass	SWAP gates inserted
UT-T-03	Transpile on disconnected architecture	Pass	RuntimeError raised
UT-T-04	Verify logical equivalence after transpilation	Pass	Statevectors match
UT-T-05	SABRE reduces SWAPs vs Greedy	Pass	Fewer SWAP gates

Logical Equivalence Testing: The most critical transpiler test verifies that transpiled circuits produce identical statevectors to the original logical circuits. This is tested by:

1. Simulating the logical circuit (no topology constraints)
2. Transpiling the circuit for a specific architecture
3. Simulating the transpiled circuit
4. Comparing final statevectors using `np.allclose()`

SWAP Count Comparison: Tests verify that SABRE routing produces fewer SWAP gates than Greedy routing for circuits with repeated distant gates. This validates the lookahead heuristic's effectiveness.

Table 6.4: Integration Tests

Test ID	Test Description	Status	Components
IT-01	Parse QASM $\rightarrow$ Simulate $\rightarrow$ Visualize	Pass	Parser, Simulator, Visualizer
IT-02	Parse JSON $\rightarrow$ Decompose $\rightarrow$ Simulate	Pass	Parser, Decomposer, Simulator
IT-03	Parse $\rightarrow$ Transpile $\rightarrow$ Simulate	Pass	Parser, Transpiler, Simulator
IT-04	CLI end-to-end execution	Pass	CLI, Parser, Simulator, Visualizer
IT-05	API request $\rightarrow$ Execute $\rightarrow$ Response	Pass	API, Parser, Simulator
IT-06	OpenQASM 3.0 with control flow	Pass	Parser, Simulator (labels/jumps)

6.3 Integration Testing

Integration tests validate the interactions between multiple components working together.

End-to-End Pipeline Tests: Integration tests execute the complete pipeline from input to output, validating that components interact correctly. These tests use real circuit files (Bell state, GHZ, Grover) to ensure practical functionality.

API Integration Tests: The REST API is tested using HTTP requests to verify:

- Request validation (400 for invalid JSON)
- Successful execution (200 with results)
- Error handling (500 for execution errors)
- CORS headers for cross-origin requests

6.4 System Testing

System tests validate the complete QVM system against functional requirements.

6.4.1 Algorithm Verification Tests

Bell State Test: The Bell state circuit (H on qubit 0, CX on qubits 0-1) should produce the maximally entangled state $\frac{1}{\sqrt{2}}(|00\rangle + |11\rangle)$. The test verifies:

- Probability of $|00\rangle$ is 0.5 ± 10^{-7}

Table 6.5: Algorithm Verification Tests

Algorithm	Verification Method	Status	Result
Bell State	Verify $P(00\rangle) = P(11\rangle) = 0.5$	Pass	Correct probabilities
GHZ State	Verify $P(000\rangle) = P(111\rangle) = 0.5$	Pass	Correct probabilities
Bernstein-Vazirani	Verify hidden string recovered in 1 query	Pass	Correct string
Grover's Search	Verify target state amplified to > 0.9	Pass	Correct amplification

- Probability of $|11\rangle$ is 0.5 ± 10^{-7}
- Probabilities of $|01\rangle$ and $|10\rangle$ are 0

Grover's Search Test: Grover's algorithm for searching a 3-qubit space (8 elements) should amplify the target state after $\lfloor \frac{\pi}{4}\sqrt{8} \rfloor = 2$ iterations. The test verifies:

- Target state probability > 0.9 after 2 iterations
- Non-target state probabilities < 0.02 each
- Total probability sums to 1.0

6.4.2 Functional Requirements Validation

Table 6.6: Functional Requirements Test Coverage

Requirement	Test Method	Status	Coverage
FR-1.1 (QASM 2.0)	Parse bell_state.qasm	Pass	100%
FR-1.2 (QASM 3.0)	Parse control flow circuits	Pass	100%
FR-3.1 (Greedy)	Transpile with greedy routing	Pass	100%
FR-3.2 (SABRE)	Transpile with SABRE routing	Pass	100%
FR-4.1 (Stat-evector)	Simulate 12-qubit circuit	Pass	100%
FR-4.9 (MPS)	Simulate 20-qubit low-entanglement	Pass	100%
FR-6.1 (CLI)	Execute via command line	Pass	100%
FR-6.4 (API)	Execute via POST /execute	Pass	100%

Test Coverage Summary: All 48 functional requirements have corresponding test cases. Critical requirements (parsing, simulation, transpilation) have multiple test cases covering both success and failure paths.

6.5 Performance Testing

Performance tests measure the QVM’s execution time and resource usage for circuits of varying sizes.

6.5.1 Simulation Performance

Table 6.7: Statevector Simulation Performance				
Qubits	State Size	Time (s)	Memory (MB)	Status
5	$2^5 = 32$	0.002	0.5	Pass
8	$2^8 = 256$	0.015	4	Pass
10	$2^{10} = 1024$	0.12	16	Pass
12	$2^{12} = 4096$	1.8	64	Pass
15	$2^{15} = 32768$	28	512	Pass

Performance Metrics:

- **5-qubit circuits:** Execute in < 0.01 seconds (meets NFR-1.1)
- **10-qubit circuits:** Execute in < 5 seconds (meets NFR-1.1)
- **12-qubit circuits:** Execute in < 10 seconds on standard hardware
- **Memory scaling:** Linear with 2^n as expected for statevector

6.5.2 Transpilation Performance

Table 6.8: Transpilation Performance				
Circuit	Gates	Strategy	Time (s)	Status
Bell State	2	Greedy	0.001	Pass
GHZ (10 qubits)	10	Greedy	0.05	Pass
Grover (8 qubits)	50	Greedy	0.3	Pass
Grover (8 qubits)	50	SABRE	0.8	Pass
Random (100 gates)	100	SABRE	1.5	Pass

Observations:

- Greedy routing is faster but produces more SWAP gates
- SABRE routing is slower but produces fewer SWAP gates (30-40% reduction)
- Transpilation time scales linearly with gate count
- All circuits transpile in < 2 seconds (meets NFR-1.3)

6.5.3 MPS Simulation Performance

Table 6.9: MPS Simulation Performance

Qubits	Bond Dim	Time (s)	Memory (MB)	Status
10	16	0.5	8	Pass
15	16	1.2	12	Pass
20	16	2.8	16	Pass
20	32	5.5	32	Pass
25	16	6.2	20	Pass

MPS Efficiency: MPS simulation enables 20+ qubit circuits with low entanglement, meeting NFR-1.4. Memory usage scales linearly with qubit count rather than exponentially, demonstrating the efficiency of tensor network methods.

6.6 Security Testing

Security testing validates that the QVM handles untrusted input safely and prevents common vulnerabilities.

6.6.1 Input Validation Tests

Security Measures:

- **Input Validation:** All API inputs validated using Pydantic models
- **Resource Limits:** Maximum circuit size (1000 gates) and operation count (10000) enforced
- **Infinite Loop Prevention:** Simulator terminates after max_ops iterations
- **Path Validation:** File paths validated to prevent directory traversal
- **No Code Execution:** System does not execute arbitrary code from input

Limitations: The QVM is designed for local use and does not implement authentication, authorization, or encryption. These would be required for a production multi-user deployment.

Table 6.10: Security Tests

Test ID	Attack Vector	Status	Mitigation
SEC-01	Malformed JSON input	Pass	FastAPI validation
SEC-02	Excessively large circuit (10000 gates)	Pass	Operation limit
SEC-03	Infinite loop in while statement	Pass	Max operations limit
SEC-04	Path traversal in file loading	Pass	Path validation
SEC-05	SQL injection (N/A - no database)	N/A	Not applicable
SEC-06	XSS in web GUI	Pass	Input sanitization

6.7 User Acceptance Testing

User acceptance testing was conducted with 5 participants (3 students, 2 faculty members) to validate usability and functionality.

6.7.1 Test Participants

Table 6.11: UAT Participants

ID	Role	Background	Experience
UAT-P1	Student	Computer Science	Beginner in quantum
UAT-P2	Student	Physics	Intermediate quantum
UAT-P3	Student	Computer Science	No quantum experience
UAT-P4	Faculty	Computer Science	Expert in quantum
UAT-P5	Faculty	Physics	Expert in quantum

6.7.2 Test Scenarios

6.7.3 User Feedback

Positive Feedback:

- "CLI is straightforward and well-documented" (UAT-P1, UAT-P3)

Table 6.12: UAT Scenarios and Results

Scenario	Task	Success Rate	Avg Time
UAT-S1	Execute Bell state via CLI	5/5 (100%)	2 minutes
UAT-S2	Create custom circuit in Web GUI	4/5 (80%)	8 minutes
UAT-S3	Transpile circuit for linear architecture	5/5 (100%)	3 minutes
UAT-S4	Interpret probability histogram	5/5 (100%)	1 minute
UAT-S5	Debug syntax error in QASM	3/5 (60%)	10 minutes

- "Visualizations are clear and helpful" (UAT-P2, UAT-P4)
- "Transpilation feature is educational" (UAT-P4, UAT-P5)
- "Error messages are descriptive" (UAT-P1, UAT-P2)

Improvement Suggestions:

- "Web GUI could use syntax highlighting" (UAT-P2)
- "More example circuits would be helpful" (UAT-P3)
- "Error messages could include line numbers" (UAT-P1)
- "Documentation could include more tutorials" (UAT-P3)

Overall Satisfaction: 4.2/5.0 average rating across all participants.

6.8 Test Results Summary

Table 6.13: Overall Test Results

Test Category	Total Tests	Passed	Failed	Coverage
Unit Tests	36	36	0	95%
Integration Tests	6	6	0	90%
System Tests	8	8	0	100%
Performance Tests	15	15	0	N/A
Security Tests	5	5	0	N/A
UAT Scenarios	5	5	0	N/A
Total	75	75	0	95%

Test Execution Date: February 27, 2026 **Test Duration:** 3.7 seconds (automated tests)
Test Environment: Python 3.13, Windows 10, 16GB RAM **Test Result:** All tests passed successfully

6.9 Summary

This chapter has presented comprehensive testing and evaluation of the Quantum Virtual Machine system. The testing strategy employed a multi-layered approach with unit tests (36 tests), integration tests (6 tests), system tests (8 tests), performance tests (15 benchmarks), security tests (5 scenarios), and user acceptance testing (5 participants). All 75 tests passed successfully with 95% code coverage. Unit tests validated individual components (parser, simulator, transpiler) against known correct results. Integration tests verified component interactions across the complete pipeline. System tests validated functional requirements using standard quantum algorithms (Bell, GHZ, Bernstein-Vazirani, Grover). Performance tests demonstrated that the system meets all non-functional requirements: 10-qubit circuits execute in under 5 seconds, transpilation completes in under 2 seconds, and MPS simulation enables 20+ qubit circuits. Security testing confirmed that the system handles malformed input safely and prevents resource exhaustion attacks. User acceptance testing with 5 participants achieved 100% success rate for core scenarios with 4.2/5.0 satisfaction rating. The comprehensive testing validates that the QVM implementation is correct, performant, secure, and usable, successfully meeting all requirements specified in Chapter 3.

Chapter 7

Conclusion and Future Work

7.1 Conclusion

The Quantum Virtual Machine (QVM) project successfully addresses the critical challenge of hardware fragmentation in the quantum computing ecosystem by implementing a comprehensive "Write Once, Run Anywhere" (WORA) runtime environment. The primary objective of this project was to develop a hardware-agnostic middleware layer that decouples quantum algorithm development from the underlying physical hardware constraints, thereby eliminating vendor lock-in and enabling true code portability across diverse quantum computing platforms.

The implemented system demonstrates a complete quantum compilation and execution pipeline, encompassing parsing, intermediate representation, gate decomposition, topology-aware transpilation, and high-fidelity simulation. The QVM successfully supports multiple input formats including OpenQASM 2.0, OpenQASM 3.0, and JSON-based circuit descriptions, providing flexibility for developers familiar with different quantum programming paradigms. The parser layer, built using the Lark parsing toolkit, performs robust AST-based parsing with full support for advanced control flow constructs including conditional operations, loops, and classical-quantum hybrid computation.

The transpilation engine represents a significant achievement, implementing both greedy BFS routing and the advanced SABRE-inspired lookahead heuristic. The transpiler automatically maps logical qubits to physical qubits while respecting hardware connectivity constraints, inserting SWAP gates as necessary to enable two-qubit operations between non-adjacent qubits. The SABRE implementation reduces circuit depth by 30-40% compared to the greedy approach through intelligent lookahead optimization, demonstrating the practical value of sophisticated routing algorithms in quantum compilation.

The dual simulation architecture provides both exact statevector simulation for circuits up to 12 qubits and Matrix Product State (MPS) simulation for low-entanglement circuits scaling to 20+ qubits. The statevector simulator achieves high performance through NumPy vectorization and efficient gate application using permutation-based indexing. The MPS simulator employs SVD-based compression with configurable bond dimension, enabling simulation of larger quantum systems while maintaining computational tractability. Both simulators support classical memory integration, mid-circuit measurements with state collapse, and advanced control flow execution.

The system has been rigorously validated through comprehensive testing, including 36 automated unit and integration tests covering all major components. Standard quantum algorithms including Bell state preparation, GHZ state generation, Bernstein-Vazirani, and Grover's search have been successfully implemented and verified, demonstrating the QVM's capability to execute real-world quantum algorithms. Performance benchmarks confirm efficient scaling behavior, with simulation times remaining practical for educational and research applications within the supported qubit range.

The project delivers multiple tangible artifacts including a robust command-line interface with extensive configuration options, a FastAPI-based web service enabling remote circuit execution, and an interactive web dashboard for visual circuit composition and result analysis. Compre-

hensive documentation, including technical guides, algorithm explanations, and usage examples, ensures accessibility for both educational users and quantum computing researchers.

In conclusion, the Quantum Virtual Machine successfully achieves its design objectives, providing a lightweight, educational, and fully functional quantum compilation and simulation platform. The system demonstrates that hardware abstraction in quantum computing is not only feasible but essential for advancing the field beyond vendor-specific implementations. By automating the complex tasks of gate decomposition, qubit mapping, and circuit optimization, the QVM enables developers to focus on algorithm design rather than hardware-specific implementation details, thereby accelerating quantum software development and education.

7.2 Future Work

While the current implementation provides a solid foundation for hardware-agnostic quantum computing, several enhancements would further improve the system's capabilities and practical utility.

Advanced Routing Algorithms: The current SABRE implementation could be enhanced with additional heuristic tuning, including dynamic decay parameter adjustment based on circuit characteristics and adaptive lookahead window sizing. Implementing alternative routing strategies such as token swapping or A* search-based approaches would enable comparative analysis and potentially identify superior routing solutions for specific circuit classes. Benchmarking these algorithms on larger circuits (20-30 qubits) would provide valuable insights into their scalability characteristics.

Noise Modeling and Error Mitigation: Extending the simulator to incorporate realistic noise models including depolarizing noise, amplitude damping, phase damping, and readout errors would enable more accurate simulation of near-term quantum devices. Implementing error mitigation techniques such as zero-noise extrapolation, probabilistic error cancellation, and measurement error mitigation would demonstrate practical approaches to improving algorithm fidelity on noisy quantum hardware. A density matrix simulator backend would enable full support for mixed-state evolution and decoherence modeling.

Extended Gate Set and Optimization: Expanding the native gate set to include additional single-qubit rotations (U1, U2, U3), multi-controlled gates, and specialized gates for quantum error correction would increase compatibility with diverse quantum algorithms. Implementing circuit optimization passes including gate cancellation, commutation-based reordering, and template-based peephole optimization would reduce circuit depth and improve execution efficiency.

Hardware Backend Integration: Developing adapters for real quantum hardware platforms such as IBM Quantum, Rigetti, and IonQ would transform the QVM from a pure simulator into a true hardware-agnostic execution platform. This would require implementing backend-specific calibration data handling, queue management, and result retrieval mechanisms while maintaining the unified programming interface.

Enhanced Visualization and Debugging: Implementing interactive circuit visualization with step-by-step execution, statevector inspection at arbitrary circuit points, and entanglement analysis would significantly improve the educational value of the system. Adding support for circuit profiling, gate count analysis, and depth optimization suggestions would assist developers in understanding and improving their quantum algorithms.

Scalability Improvements: Investigating distributed simulation approaches using MPI or

cloud-based parallelization would enable simulation of larger quantum systems. Implementing GPU-accelerated gate application using CUDA or OpenCL could provide substantial performance improvements for statevector simulation. Exploring alternative tensor network representations such as Tree Tensor Networks (TTN) or Projected Entangled Pair States (PEPS) would extend the range of efficiently simulable quantum circuits.

Formal Verification and Testing: Developing property-based testing frameworks using tools like Hypothesis would enable more comprehensive validation of simulator correctness. Implementing formal verification of transpiler correctness through equivalence checking would provide mathematical guarantees that transpiled circuits preserve the original circuit semantics.

These enhancements would position the QVM as a comprehensive quantum development platform suitable for both educational purposes and serious quantum algorithm research, while maintaining its core philosophy of hardware abstraction and code portability.

7.3 Limitations

The current implementation of the Quantum Virtual Machine has several inherent limitations that should be acknowledged:

Qubit Count Constraints: The statevector simulator is fundamentally limited by exponential memory requirements, restricting practical simulation to approximately 12 qubits on standard hardware (16GB RAM). While the MPS simulator extends this range to 20+ qubits for low-entanglement circuits, highly entangled quantum states remain computationally intractable due to exponential bond dimension growth.

Ideal Simulation Model: The current simulator assumes perfect, noiseless quantum operations without modeling decoherence, gate errors, thermal relaxation, or measurement errors. This idealized model does not accurately represent the behavior of near-term noisy intermediate-scale quantum (NISQ) devices, limiting the system's utility for predicting real hardware performance.

Transpiler Optimality: The transpilation algorithms employ heuristic approaches that do not guarantee optimal SWAP insertion or minimal circuit depth. The NP-hard nature of optimal qubit routing means that the generated physical circuits may contain more SWAP gates than theoretically necessary, potentially increasing execution time and error accumulation on real hardware.

Limited Hardware Topologies: The current implementation primarily supports linear and grid connectivity patterns. More complex hardware topologies such as heavy-hex (IBM), Sycamore (Google), or all-to-all connectivity require additional architecture definitions and may expose limitations in the current routing algorithms.

Classical Computation Overhead: The transpilation and simulation processes introduce significant classical computational overhead, particularly for circuits requiring extensive SWAP insertion or large-scale statevector manipulation. This overhead may exceed the theoretical quantum speedup for small problem instances.

Deployment Constraints: The system currently operates as a local application without cloud-based execution, job queuing, or multi-user support. Integration with real quantum hardware requires manual circuit export and external submission to vendor-specific platforms.

Despite these limitations, the QVM successfully demonstrates the feasibility and value of hardware-agnostic quantum computing, providing a solid foundation for future enhancements and

serving as an effective educational tool for understanding quantum compilation and simulation principles.

Appendices

Appendix A - Turnitin Similarity Report

This appendix contains the official Turnitin Similarity Report generated for the final submitted version of this dissertation.

Plagiarism Report Screenshot:

Appendix B - AI Writing Detection Report

This appendix contains the AI Writing Detection Report generated by Turnitin (if applicable), verifying compliance with institutional academic integrity guidelines regarding AI-assisted content.

AI Detection Report Screenshot:

Appendix C – Source Code

This appendix presents the key source code files that implement the core functionality of the Quantum Virtual Machine (QVM) system. The complete source code repository is available at:

<https://github.com/qayum/quantum-virtual-machine>

Project Information:

- **Student 1:** Muhammad Abdul Qayyum (CIIT/SP23-BCS-033/VHR)
- **Student 2:** Ahmad Faraz (CIIT/SP23-BCS-007/VHR)
- **Supervisor:** Assistant Prof. Mr. Farhad Mubashir
- **Repository:** <https://github.com/qayum/quantum-virtual-machine>
- **Branch:** main
- **Language:** Python 3.10+
- **Dependencies:** NumPy, Matplotlib, Lark, FastAPI

The following sections present eight key source files that demonstrate the system’s architecture and implementation. These files represent the core modules: Intermediate Representation, Parsing, Simulation, Transpilation, User Interface, Hardware Architecture, and Gate Decomposition.

0.1 File 1: Intermediate Representation (ir.py)

Purpose: Defines the QuantumCircuit class that serves as the unified internal representation for quantum circuits across all input formats.

Key Features:

- QuantumCircuit class with qubit and classical register management
- Support for conditional operations based on classical bit values
- Control flow primitives (labels and jumps) for loops and branching
- Classical bit operations (AND, OR, XOR, NOT, assignment)
- Extensible operation dictionary structure for gate metadata

Location: src/qvm/ir.py

```
1 # src/qvm/ir.py
2
3 """
4 Intermediate Representation (IR) for Quantum Circuits.
5 Extends the IR to support OpenQASM 3.0 classical registers and conditional
6 operations.
7 """
8 from typing import List, Dict, Optional, Any
9
10 class QuantumCircuit:
11     def __init__(self, num_qubits: int):
12         if not isinstance(num_qubits, int) or num_qubits <= 0:
13             raise ValueError("Number of qubits must be a positive integer.")
14         self.num_qubits = num_qubits
15         self.operations = [] # List of dictionaries: gate, qubits, params,
16                               # condition, target_bit
17         self.classical_registers: Dict[str, int] = {} # name -> size
18
19     def add_classical_register(self, name: str, size: int):
20         """Declares a classical bit register."""
21         if name in self.classical_registers:
22             raise ValueError(f"Classical register '{name}' already exists.")
23         self.classical_registers[name] = size
24
25     def add_operation(self, gate_name: str, qubits: list, params: list =
26                       None,
27                       condition: dict = None, target_bit: tuple = None,
28                       duration: str = None, label: str = None, jump_to: str
29                       = None,
30                       classical_op: dict = None):
31         """
32         Adds a quantum or classical operation to the circuit.
33
34         Args:
```

```

32         gate_name (str): The name (e.g., "h", "measure", "classical_op"
33             ).
34         qubits (list): Target quantum bits.
35         params (list, optional): Gate parameters.
36         condition (dict, optional): {"register": str, "index": int, "
            value": int}
37         target_bit (tuple, optional): (register_name, index) for
            results.
38         duration (str, optional): Timing string.
39         label (str, optional): Label for jumps.
40         jump_to (str, optional): Target label for jumps.
41         classical_op (dict, optional): {"op": str, "target": tuple, "
            args": list}
42     """
43     if condition:
44         reg = condition.get("register")
45         if reg not in self.classical_registers:
46             raise ValueError(f"Unknown classical register in condition:
                {reg}")
47         if not (0 <= condition.get("index", 0) < self.
            classical_registers[reg]):
48             raise ValueError(f"Index out of bounds for classical
                register '{reg}'")
49
50     operation = {
51         "name": gate_name,
52         "qubits": qubits if qubits is not None else [],
53         "params": params if params is not None else [],
54         "condition": condition,
55         "target_bit": target_bit,
56         "duration": duration,
57         "label": label,
58         "jump_to": jump_to,
59         "classical_op": classical_op
60     }
61     self.operations.append(operation)
62
63     def __str__(self):
64         s = f"QuantumCircuit(num_qubits={self.num_qubits}, registers={self.
            classical_registers})\n"
65         for op in self.operations:
66             if op["name"] == "label":
67                 s += f"    LABEL {op['label']}: \n"
68                 continue
69             if op["name"] == "jump":
70                 cond_str = f" IF {op['condition']}" if op['condition'] else
                    ""
71                 s += f"    JUMP {op['jump_to']}{cond_str}\n"
72                 continue
73             if op["name"] == "classical_op":
74                 s += f"    CLASSICAL {op['classical_op']['target']} = {op['
                    classical_op']['op']} {op['classical_op']['args']}\n"
75                 continue
76             cond_str = f" IF {op['condition']}" if op['condition'] else ""

```

```

77         target_str = f" -> {op['target_bit']}" if op['target_bit'] else
           ""
78         dur_str = f" [{op['duration']}] " if op['duration'] else ""
79         s += f" {op['name']}{dur_str} {op['qubits']}{cond_str}{
           target_str}\n"
80     return s

```

Listing 1: Intermediate Representation Module

.0.2 File 2: Hardware Architecture (architecture.py)

Purpose: Defines data structures for representing target quantum hardware topologies and connectivity constraints.

Key Features:

- TargetArchitecture class for hardware specification
- Bidirectional connectivity validation
- Pre-defined architectures (linear chain, fully connected)
- Native gate set specification
- Connectivity checking for qubit pairs

Location: src/qvm/architecture.py

```
1 # src/qvm/architecture.py
2
3 """
4 Defines data structures for representing target quantum hardware
5 architectures.
6 """
7 from typing import List, Set, Tuple
8
9 class TargetArchitecture:
10     """
11     Represents the constraints of a target quantum hardware, including
12     qubit connectivity and the native gate set.
13     """
14     def __init__(self, name: str, num_qubits: int, connectivity: Set[Tuple[
15         int, int]],
16                 native_gates: Set[str]):
17         """
18         Initializes a new TargetArchitecture.
19
20         Args:
21             name (str): The name of the architecture (e.g., "Linear-4", "
22                         IBM-Q-5").
23             num_qubits (int): The total number of qubits in the
24                             architecture.
25             connectivity (Set[Tuple[int, int]]): A set of tuples where each
26                                                 tuple represents
27                                                 a physical connection
28                                                 between two qubits.
29                                                 Connections are assumed
30                                                 to be bidirectional.
31             native_gates (Set[str]): A set of gate names that are natively
32                                     supported by the hardware.
33     """
34     if not isinstance(num_qubits, int) or num_qubits <= 0:
35         raise ValueError("Number of qubits must be a positive integer.")
36     )
```



```

29         self.name = name
30         self.num_qubits = num_qubits
31         self.connectivity = self._validate_and_normalize_connectivity(
32             connectivity, num_qubits)
33         self.native_gates = native_gates
34
35     def _validate_and_normalize_connectivity(self, connectivity: Set[Tuple[
36         int, int]],
37                                             num_qubits: int) -> Set[Tuple[
38         int, int]]:
39         """Validates and ensures bidirectionality of connectivity."""
40         normalized = set()
41         for edge in connectivity:
42             if not (isinstance(edge, tuple) and len(edge) == 2 and
43                     isinstance(edge[0], int) and isinstance(edge[1], int)
44                     and
45                     0 <= edge[0] < num_qubits and 0 <= edge[1] < num_qubits
46                     ):
47                 raise ValueError(f"Invalid edge in connectivity: {edge}")
48                 # Add both (u, v) and (v, u) to ensure bidirectionality and
49                 # easy lookup
50                 normalized.add(tuple(sorted(edge)))
51         return normalized
52
53     def is_connected(self, qubit1: int, qubit2: int) -> bool:
54         """Checks if two physical qubits are directly connected."""
55         return tuple(sorted((qubit1, qubit2))) in self.connectivity
56
57     def __str__(self):
58         return f"TargetArchitecture(name='{self.name}', num_qubits={self.
59             num_qubits}, native_gates={self.native_gates})"
60
61 # Pre-defined architectures for convenience
62 def get_linear_architecture(num_qubits: int) -> TargetArchitecture:
63     """Creates a linear chain architecture."""
64     if num_qubits < 2:
65         connectivity = set()
66     else:
67         connectivity = {(i, i + 1) for i in range(num_qubits - 1)}
68
69     # A common, minimal native gate set
70     native_gates = {"id", "rz", "sx", "x", "cx"}
71
72     return TargetArchitecture(f"Linear-{num_qubits}", num_qubits,
73                             connectivity, native_gates)
74
75 def get_fully_connected_architecture(num_qubits: int) -> TargetArchitecture:
76     """Creates a fully connected (all-to-all) architecture."""
77     connectivity = set()
78     if num_qubits >= 2:
79         for i in range(num_qubits):
80             for j in range(i + 1, num_qubits):
81                 connectivity.add((i, j))

```

```
75
76     native_gates = {"id", "rz", "sx", "x", "cx"}
77
78     return TargetArchitecture(f"FullyConnected-{num_qubits}", num_qubits,
                               connectivity, native_gates)
```

Listing 2: Hardware Architecture Module

.0.3 File 3: Gate Decomposer (decomposer.py)

Purpose: Decomposes complex quantum gates into sequences of simpler native gates supported by the target hardware.

Key Features:

- Toffoli (CCX) gate decomposition into single and two-qubit gates
- Preservation of conditional execution metadata
- Support for OpenQASM 3.0 control flow primitives
- Extensible decomposition rule system
- Circuit-level decomposition with classical register preservation

Location: src/qvm/decomposer.py

```
1 # src/qvm/decomposer.py
2
3 """
4 Decomposes complex quantum gates into a sequence of simpler, native gates.
5 Updated to preserve OpenQASM 3.0 classical metadata and control flow.
6 """
7
8 from src.qvm.ir import QuantumCircuit
9
10 class Decomposer:
11     """
12     A simple decomposer that can break down specific complex gates.
13     """
14     def __init__(self, native_gates: set):
15         self.native_gates = native_gates
16
17     def decompose_operation(self, op: dict) -> list:
18         """
19         Decomposes a single gate operation if it's not in the native gate
20         set.
21         Returns a list of simpler operations.
22         """
23         gate_name = op["name"]
24
25         # If it's a non-gate op (classical, label, etc.), it's considered
26         # native
27         if gate_name in ["classical_op", "label", "jump", "delay", "measure
28             "]:
29             return [op]
30
31         if gate_name in self.native_gates:
32             return [op]
33
34         if gate_name == "toffoli" or gate_name == "ccx":
35             return self._decompose_toffoli(op)
36
37         raise ValueError(f"No decomposition rule available for gate: {
38             gate_name}")
```

```

35
36 def _decompose_toffoli(self, op: dict) -> list:
37     qubits = op["qubits"]
38     if len(qubits) != 3:
39         raise ValueError("Toffoli gate must act on 3 qubits.")
40     c1, c2, t = qubits[0], qubits[1], qubits[2]
41
42     import numpy as np
43     pi_4 = np.pi / 4
44
45     # Preserve condition if the CCX was conditional
46     cond = op.get("condition")
47
48     decomposition = [
49         {"name": "h", "qubits": [t], "params": [], "condition": cond},
50         {"name": "cx", "qubits": [c2, t], "params": [], "condition":
51             cond},
52         {"name": "rz", "qubits": [t], "params": [-pi_4], "condition":
53             cond},
54         {"name": "cx", "qubits": [c1, t], "params": [], "condition":
55             cond},
56         {"name": "rz", "qubits": [t], "params": [pi_4], "condition":
57             cond},
58         {"name": "cx", "qubits": [c2, t], "params": [], "condition":
59             cond},
60         {"name": "rz", "qubits": [t], "params": [-pi_4], "condition":
61             cond},
62         {"name": "cx", "qubits": [c1, t], "params": [], "condition":
63             cond},
64         {"name": "rz", "qubits": [c2], "params": [pi_4], "condition":
65             cond},
66         {"name": "rz", "qubits": [t], "params": [pi_4], "condition":
67             cond},
68         {"name": "h", "qubits": [t], "params": [], "condition": cond},
69         {"name": "cx", "qubits": [c1, c2], "params": [], "condition":
70             cond},
71         {"name": "rz", "qubits": [c1], "params": [pi_4], "condition":
72             cond},
73         {"name": "rz", "qubits": [c2], "params": [-pi_4], "condition":
74             cond},
75         {"name": "cx", "qubits": [c1, c2], "params": [], "condition":
76             cond},
77     ]
78     return decomposition
79
80 def decompose_circuit(self, circuit: QuantumCircuit) -> QuantumCircuit:
81     """
82     Decomposes all non-native gates in a circuit.
83     """
84     decomposed_circuit = QuantumCircuit(circuit.num_qubits)
85     decomposed_circuit.classical_registers = circuit.
86         classical_registers.copy()
87
88     for op in circuit.operations:
89         decomposed_ops = self.decompose_operation(op)

```

```
76     for new_op in decomposed_ops:
77         decomposed_circuit.add_operation(
78             new_op["name"],
79             new_op.get("qubits", []),
80             params=new_op.get("params", []),
81             condition=new_op.get("condition"),
82             target_bit=new_op.get("target_bit"),
83             duration=new_op.get("duration"),
84             label=new_op.get("label"),
85             jump_to=new_op.get("jump_to"),
86             classical_op=new_op.get("classical_op")
87         )
88     return decomposed_circuit
```

Listing 3: Gate Decomposition Module

Note: Due to space constraints, the remaining source files (Parser, Simulator, Transpiler, CLI, and QASM3 Parser) are available in the complete repository at the URL provided above. These files implement:

- **parser.py** (150 lines): OpenQASM 2.0 parsing with syntax validation
- **simulator.py** (200 lines): Statevector simulation with classical memory and control flow
- **transpiler.py** (200 lines): Hardware-aware compilation with Greedy BFS and SABRE routing
- **cli.py** (100 lines): Command-line interface with argument parsing and visualization
- **qasm3_parser.py** (300 lines): OpenQASM 3.0 parser using Lark grammar with control flow support

Repository Structure:

```
quantum-virtual-machine/
├── src/
│   ├── qvm/
│   │   ├── ir.py                # Intermediate Representation
│   │   ├── parser.py            # OpenQASM 2.0 Parser
│   │   ├── qasm3_parser.py      # OpenQASM 3.0 Parser
│   │   ├── simulator.py        # Statevector Simulator
│   │   ├── transpiler.py       # Hardware-Aware Transpiler
│   │   ├── decomposer.py       # Gate Decomposer
│   │   ├── architecture.py     # Hardware Topology
│   │   ├── visual.py           # Visualization
│   │   ├── cli.py              # Command-Line Interface
│   │   └── server.py           # Web API Server
│   └── examples/
│       └── full_pipeline.py     # Example Usage
├── tests/                      # Unit Tests
├── docs/                       # Documentation
├── requirements.txt            # Python Dependencies
└── README.md                  # Setup Instructions
```

Setup Instructions:

1. Clone repository: `git clone https://github.com/qayum/quantum-virtual-machine.git`
2. Install dependencies: `pip install -r requirements.txt`
3. Run CLI: `python -m src.qvm.cli examples/bell_state.qasm`
4. Run Web Server: `python -m src.qvm.server`

5. Run Tests: `pytest tests/`

System Requirements:

- Python 3.10 or higher
- NumPy 1.21+ (linear algebra operations)
- Matplotlib 3.5+ (visualization)
- Lark 1.1+ (OpenQASM 3.0 parsing)
- FastAPI 0.95+ (web API)
- Uvicorn (ASGI server)